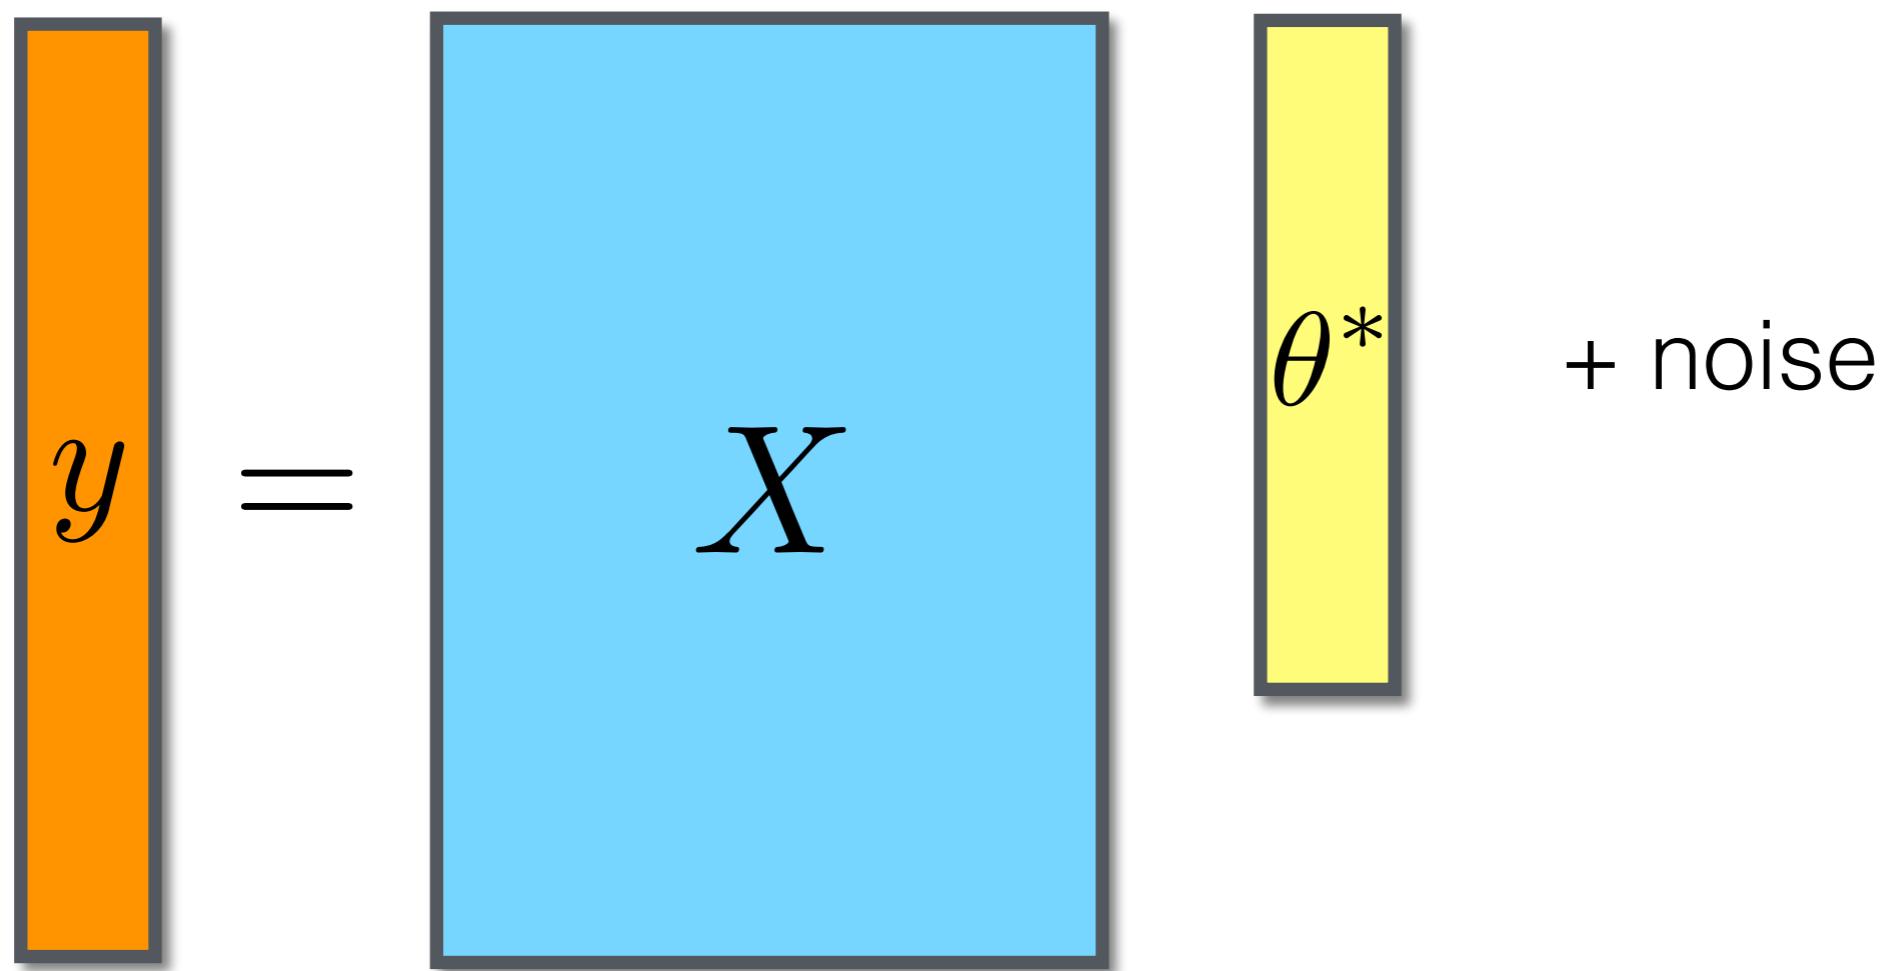


Linear Regression with Graph Constraints

Ludwig Schmidt
MIT → UC Berkeley

Based on joint works with Tobias Ehrenberger,
Ernest Fraenkel, Chinmay Hegde, Piotr Indyk, Stefanie
Jegelka, Alexander LeNail, Jonathan Li, and Karen Sachs.

Linear Regression

$$y = X \theta^* + \text{noise}$$


Given: data matrix X and observations y .

Goal: find the unknown θ^* .
(also: predict y in the future)

Linear Regression

$$y = X \theta^* + \text{noise}$$

Given: data matrix X and observations y .

Goal: find the unknown θ^* .
(also: predict y in the future)



Gauss, early 1800s

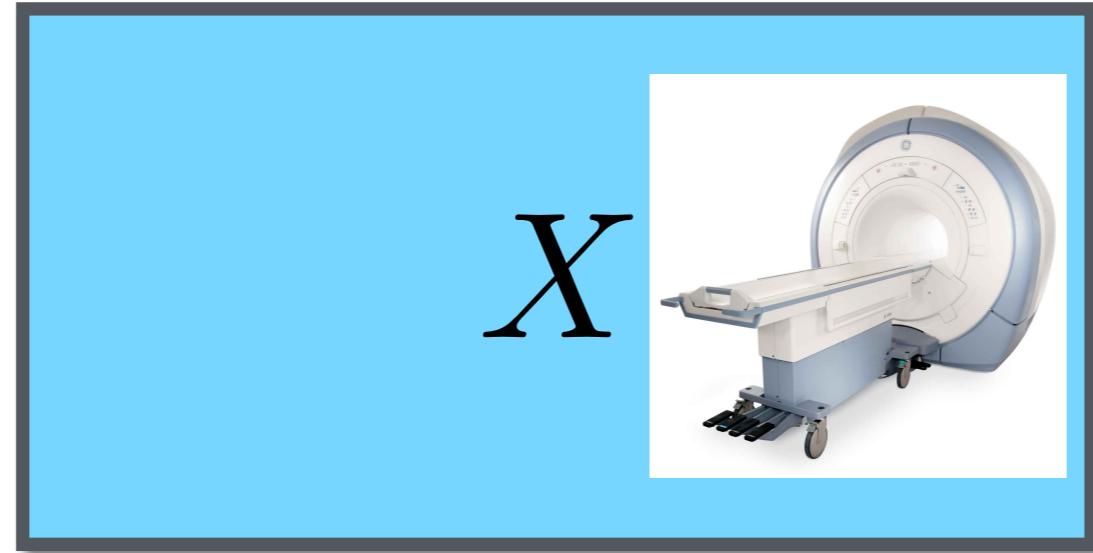
Constrained Linear Regression

$$n \begin{matrix} y \\ \end{matrix} = \begin{matrix} d \\ X \\ \end{matrix} \theta^*$$

In general, **impossible** if $n < d$.

But: many applications (MRI, GWAS, ...)

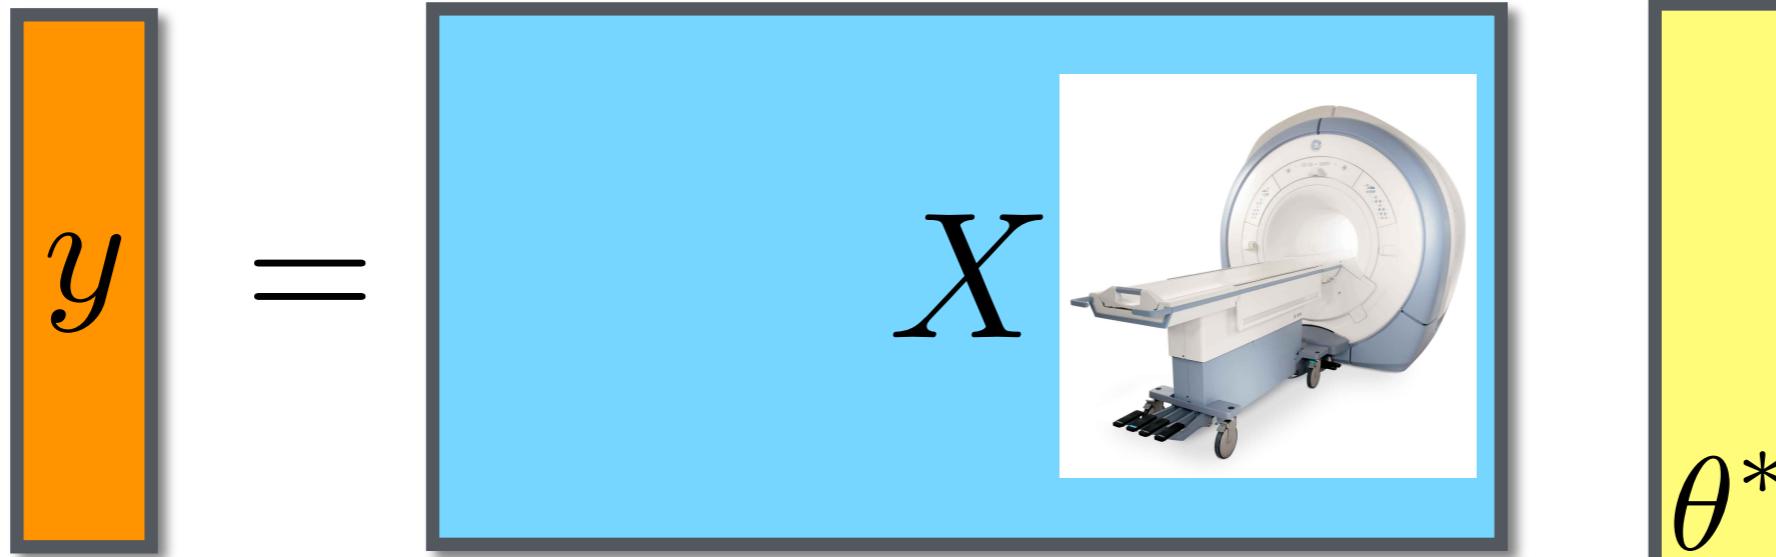
Constrained Linear Regression

$$n \begin{array}{|c|} \hline y \\ \hline \end{array} = \begin{array}{|c|} \hline d \\ \hline X \\ \hline \end{array} \begin{array}{|c|} \hline \theta^* \\ \hline \end{array}$$


In general, **impossible** if $n < d$.

But: many applications (MRI, GWAS, ...)

Constrained Linear Regression

$$n \begin{matrix} y \\ \end{matrix} = \begin{matrix} d \\ X \end{matrix} \theta^*$$


In general, **impossible** if $n < d$.

But: many applications (MRI, GWAS, ...)



→ Use extra structure such as **sparsity**.

Solving Linear Regression with Constraints

Standard
Linear Regression

Sparse
Linear Regression

Algorithm

$$\arg \min_{\theta} \|y - X\theta\|_2^2$$

Least Squares

Solving Linear Regression with Constraints

Standard
Linear Regression

Algorithm

$$\arg \min_{\theta} \|y - X\theta\|_2^2$$

Least Squares

Sparse
Linear Regression

$$\arg \min_{\theta} \|y - X\theta\|_2^2$$

$$\|\theta\|_0 \leq \lambda$$

Solving Linear Regression with Constraints

Standard
Linear Regression

Algorithm

$$\arg \min_{\theta} \|y - X\theta\|_2^2$$

Least Squares

Sparse
Linear Regression

$$\begin{aligned} & \arg \min \|y - X\theta\|_2^2 \\ & \|\theta\|_1 \leq \lambda \end{aligned}$$

Lasso

Convex relaxation $\ell_0 \rightarrow \ell_1$

Solving Linear Regression with Constraints

Standard
Linear Regression

Sparse
Linear Regression

Algorithm

$$\arg \min_{\theta} \|y - X\theta\|_2^2$$

Least Squares

$$\arg \min_{\theta} \|y - X\theta\|_2^2 \\ \|\theta\|_1 \leq \lambda$$

Lasso

Convex relaxation $\ell_0 \rightarrow \ell_1$

Sample
Complexity

$$n = O(d)$$

$$n = O(s \log d)$$

(d is the ambient dimension, s is the sparsity of the unknown vector)

Widely Used Methods



Rob Tibshirani

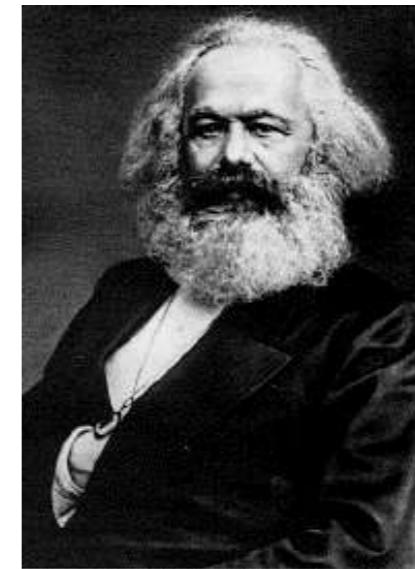
Popularized the Lasso

Widely Used Methods



Rob Tibshirani

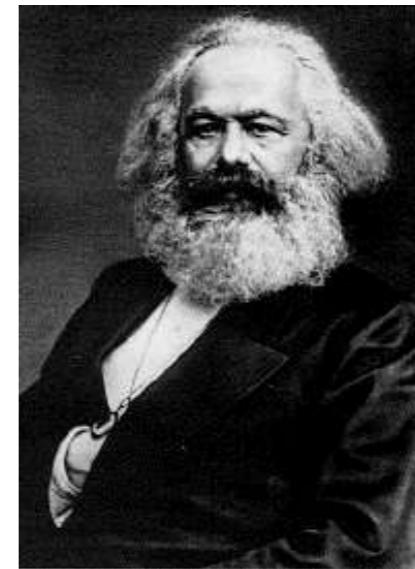
Popularized the Lasso



Karl Marx

Popularized Communism

Widely Used Methods



Rob Tibshirani

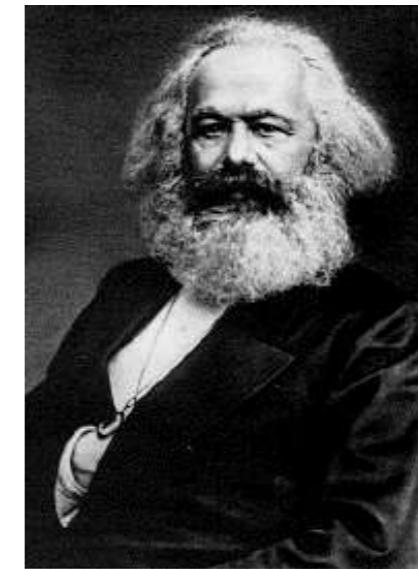
Karl Marx

Popularized the Lasso

Popularized Communism

What do they have in common?

Widely Used Methods



Rob Tibshirani

Popularized the Lasso

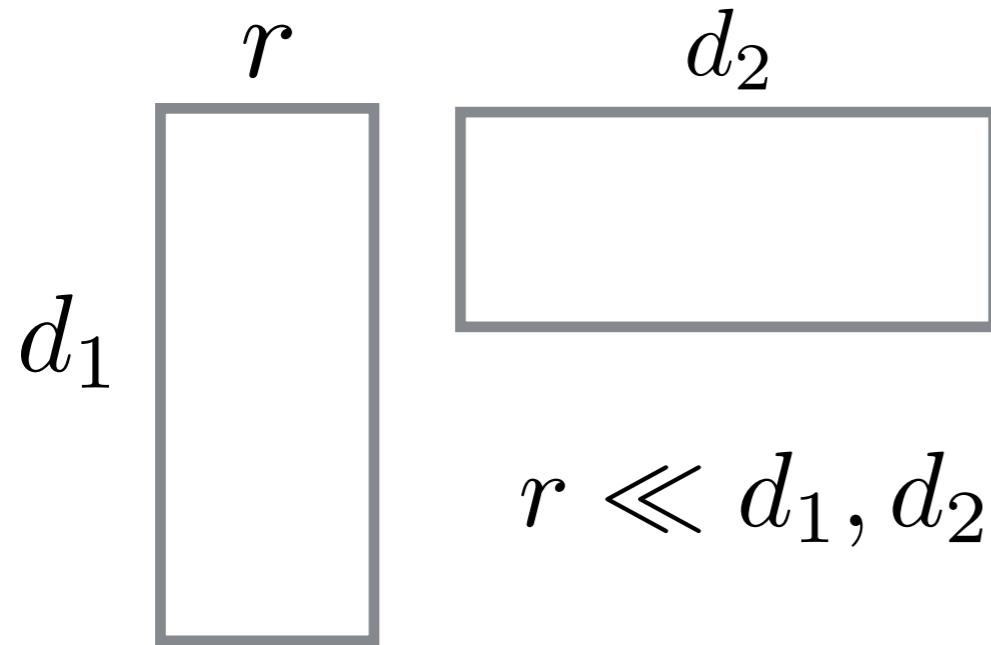
Google Scholar:
266,534 citations

Karl Marx

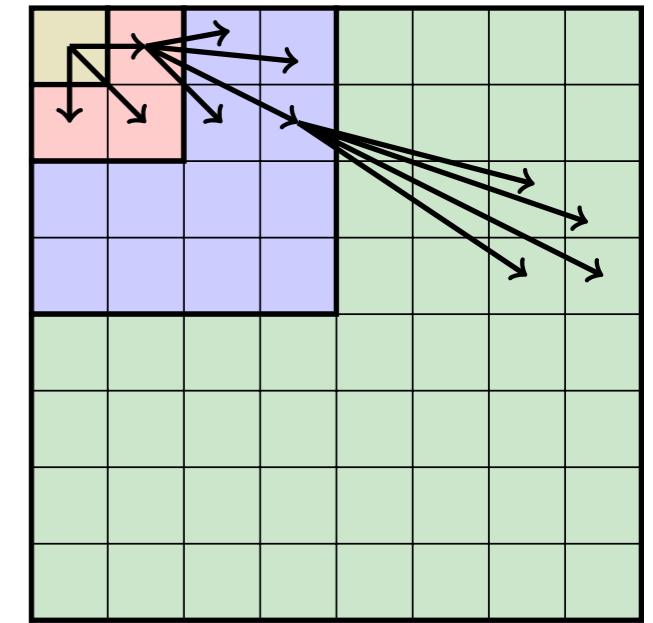
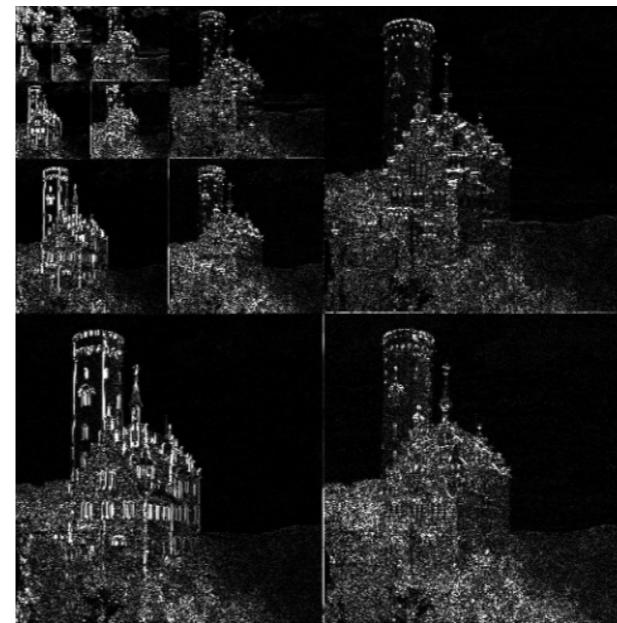
Popularized Communism

Google Scholar:
275,099 citations

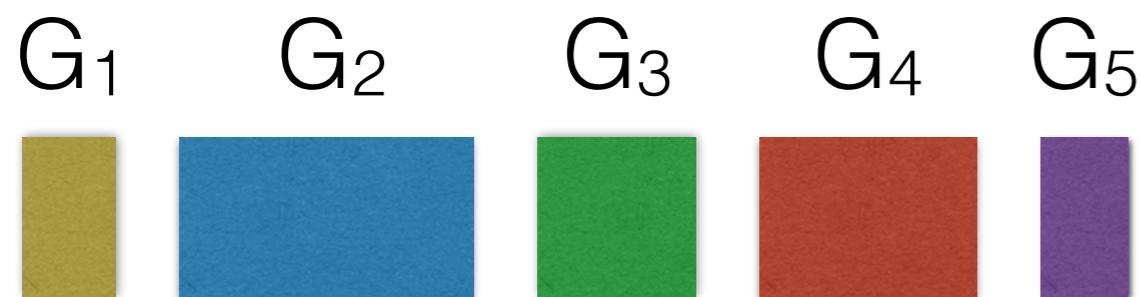
Wide Range of Constraints



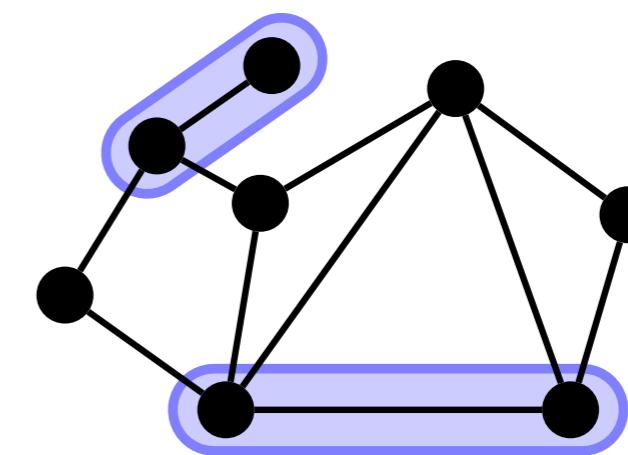
Low rank



Tree sparsity

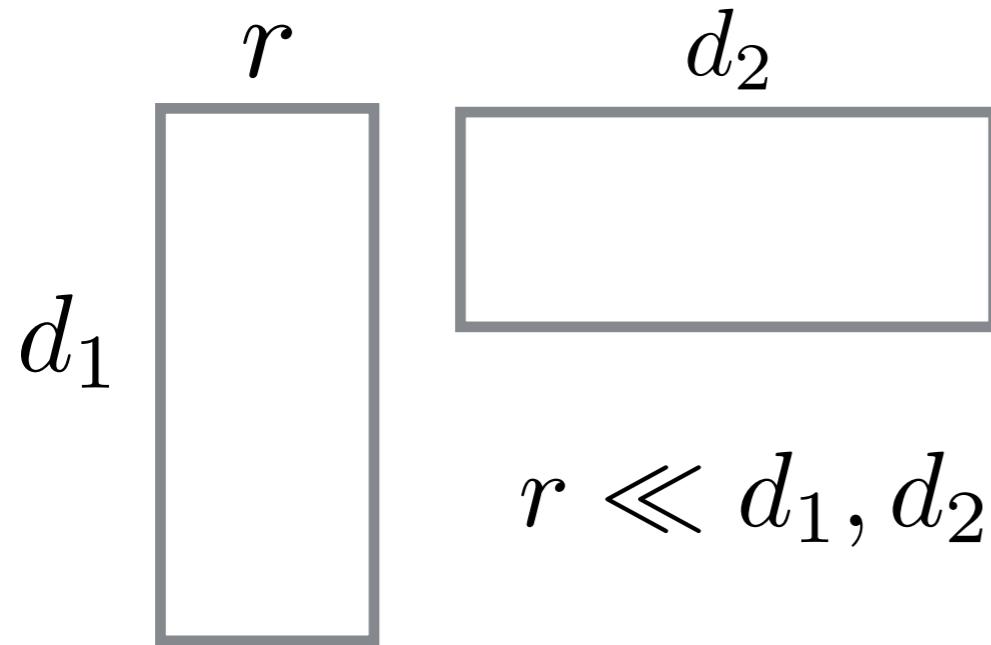


Group sparsity

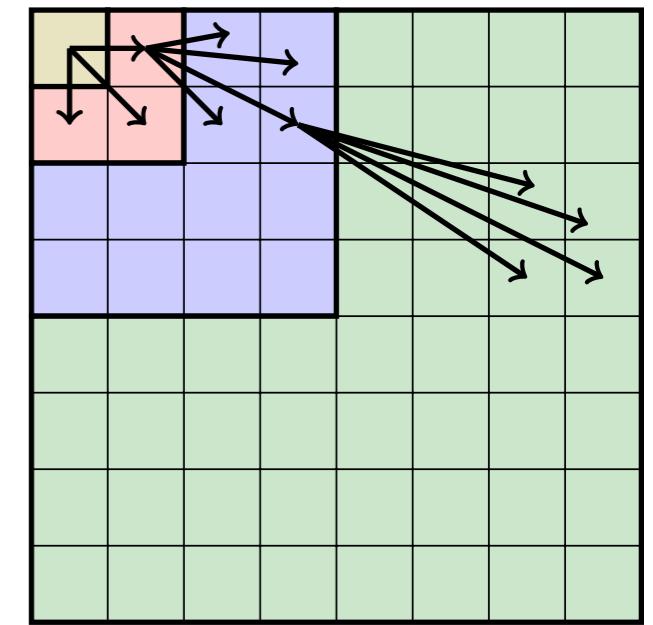
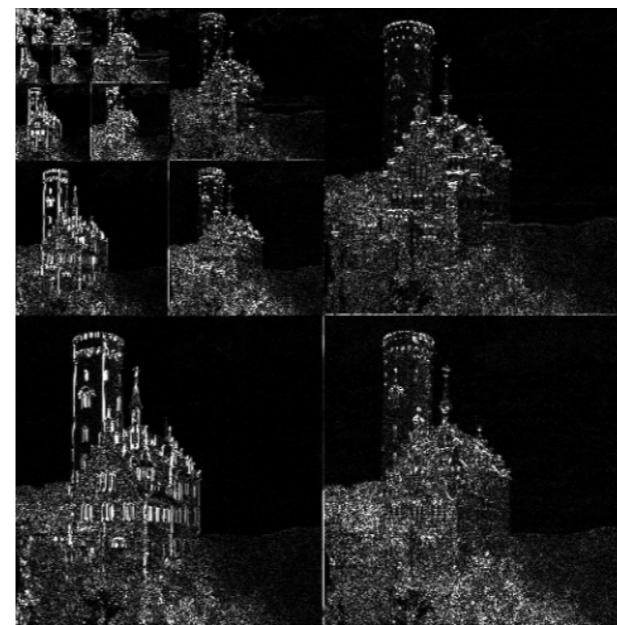


Graph sparsity

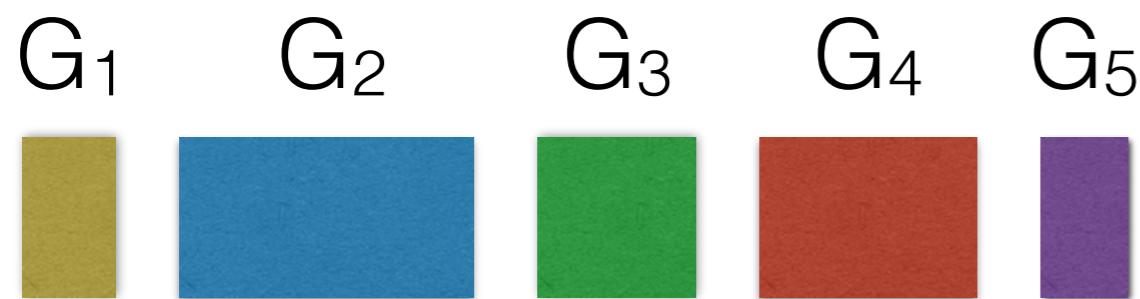
Wide Range of Constraints



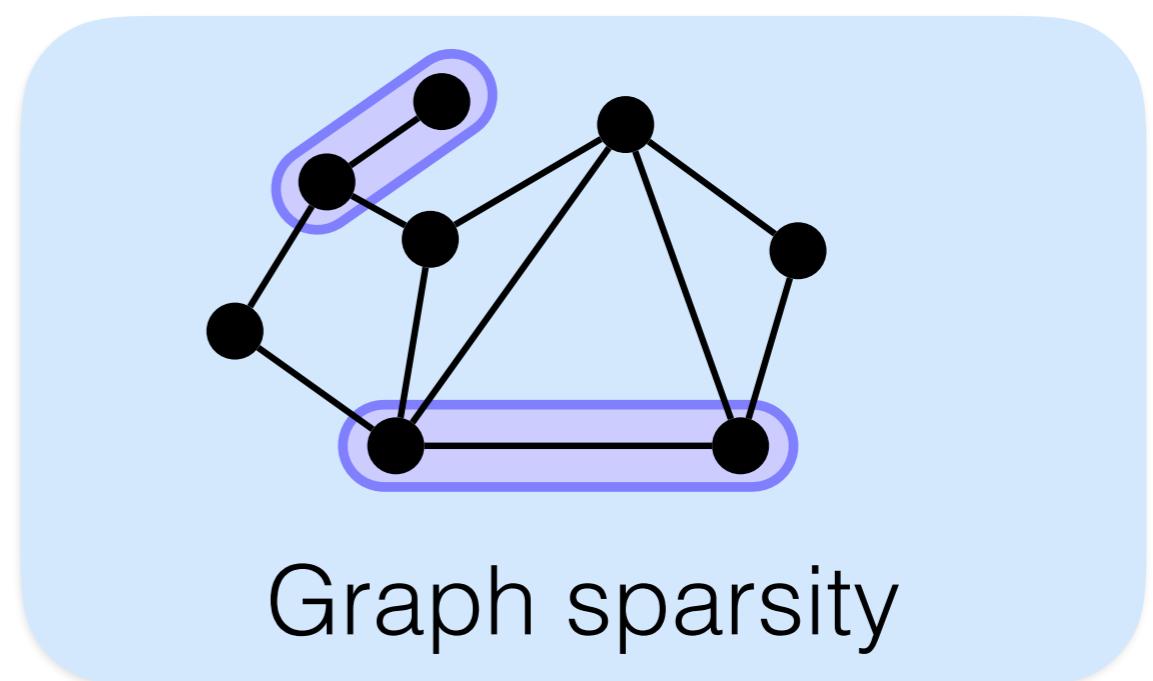
Low rank



Tree sparsity

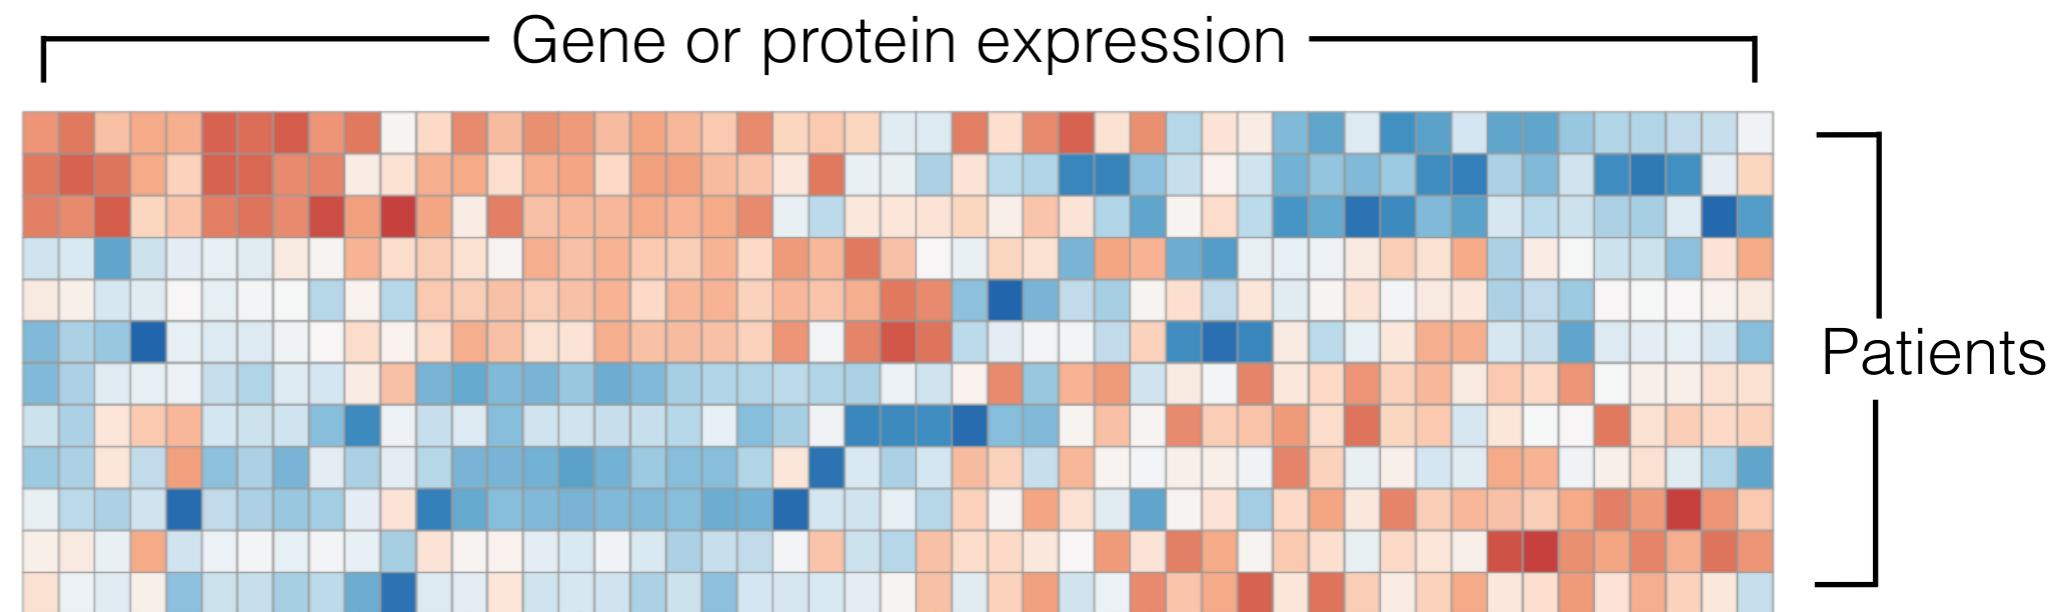


Group sparsity

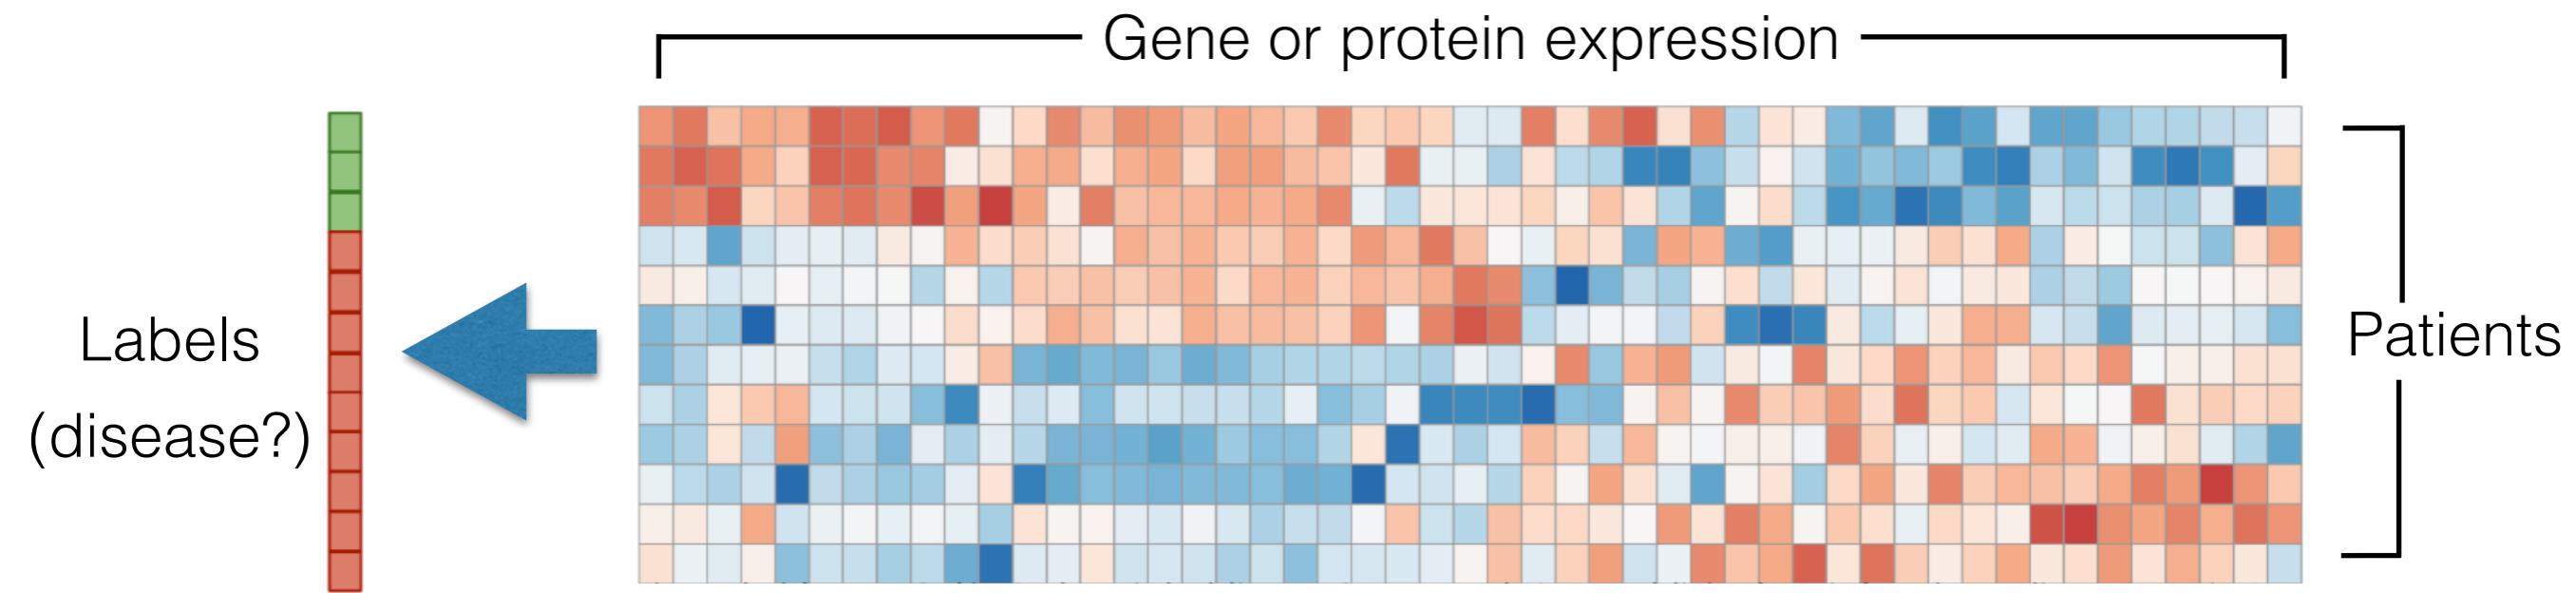


Graph sparsity

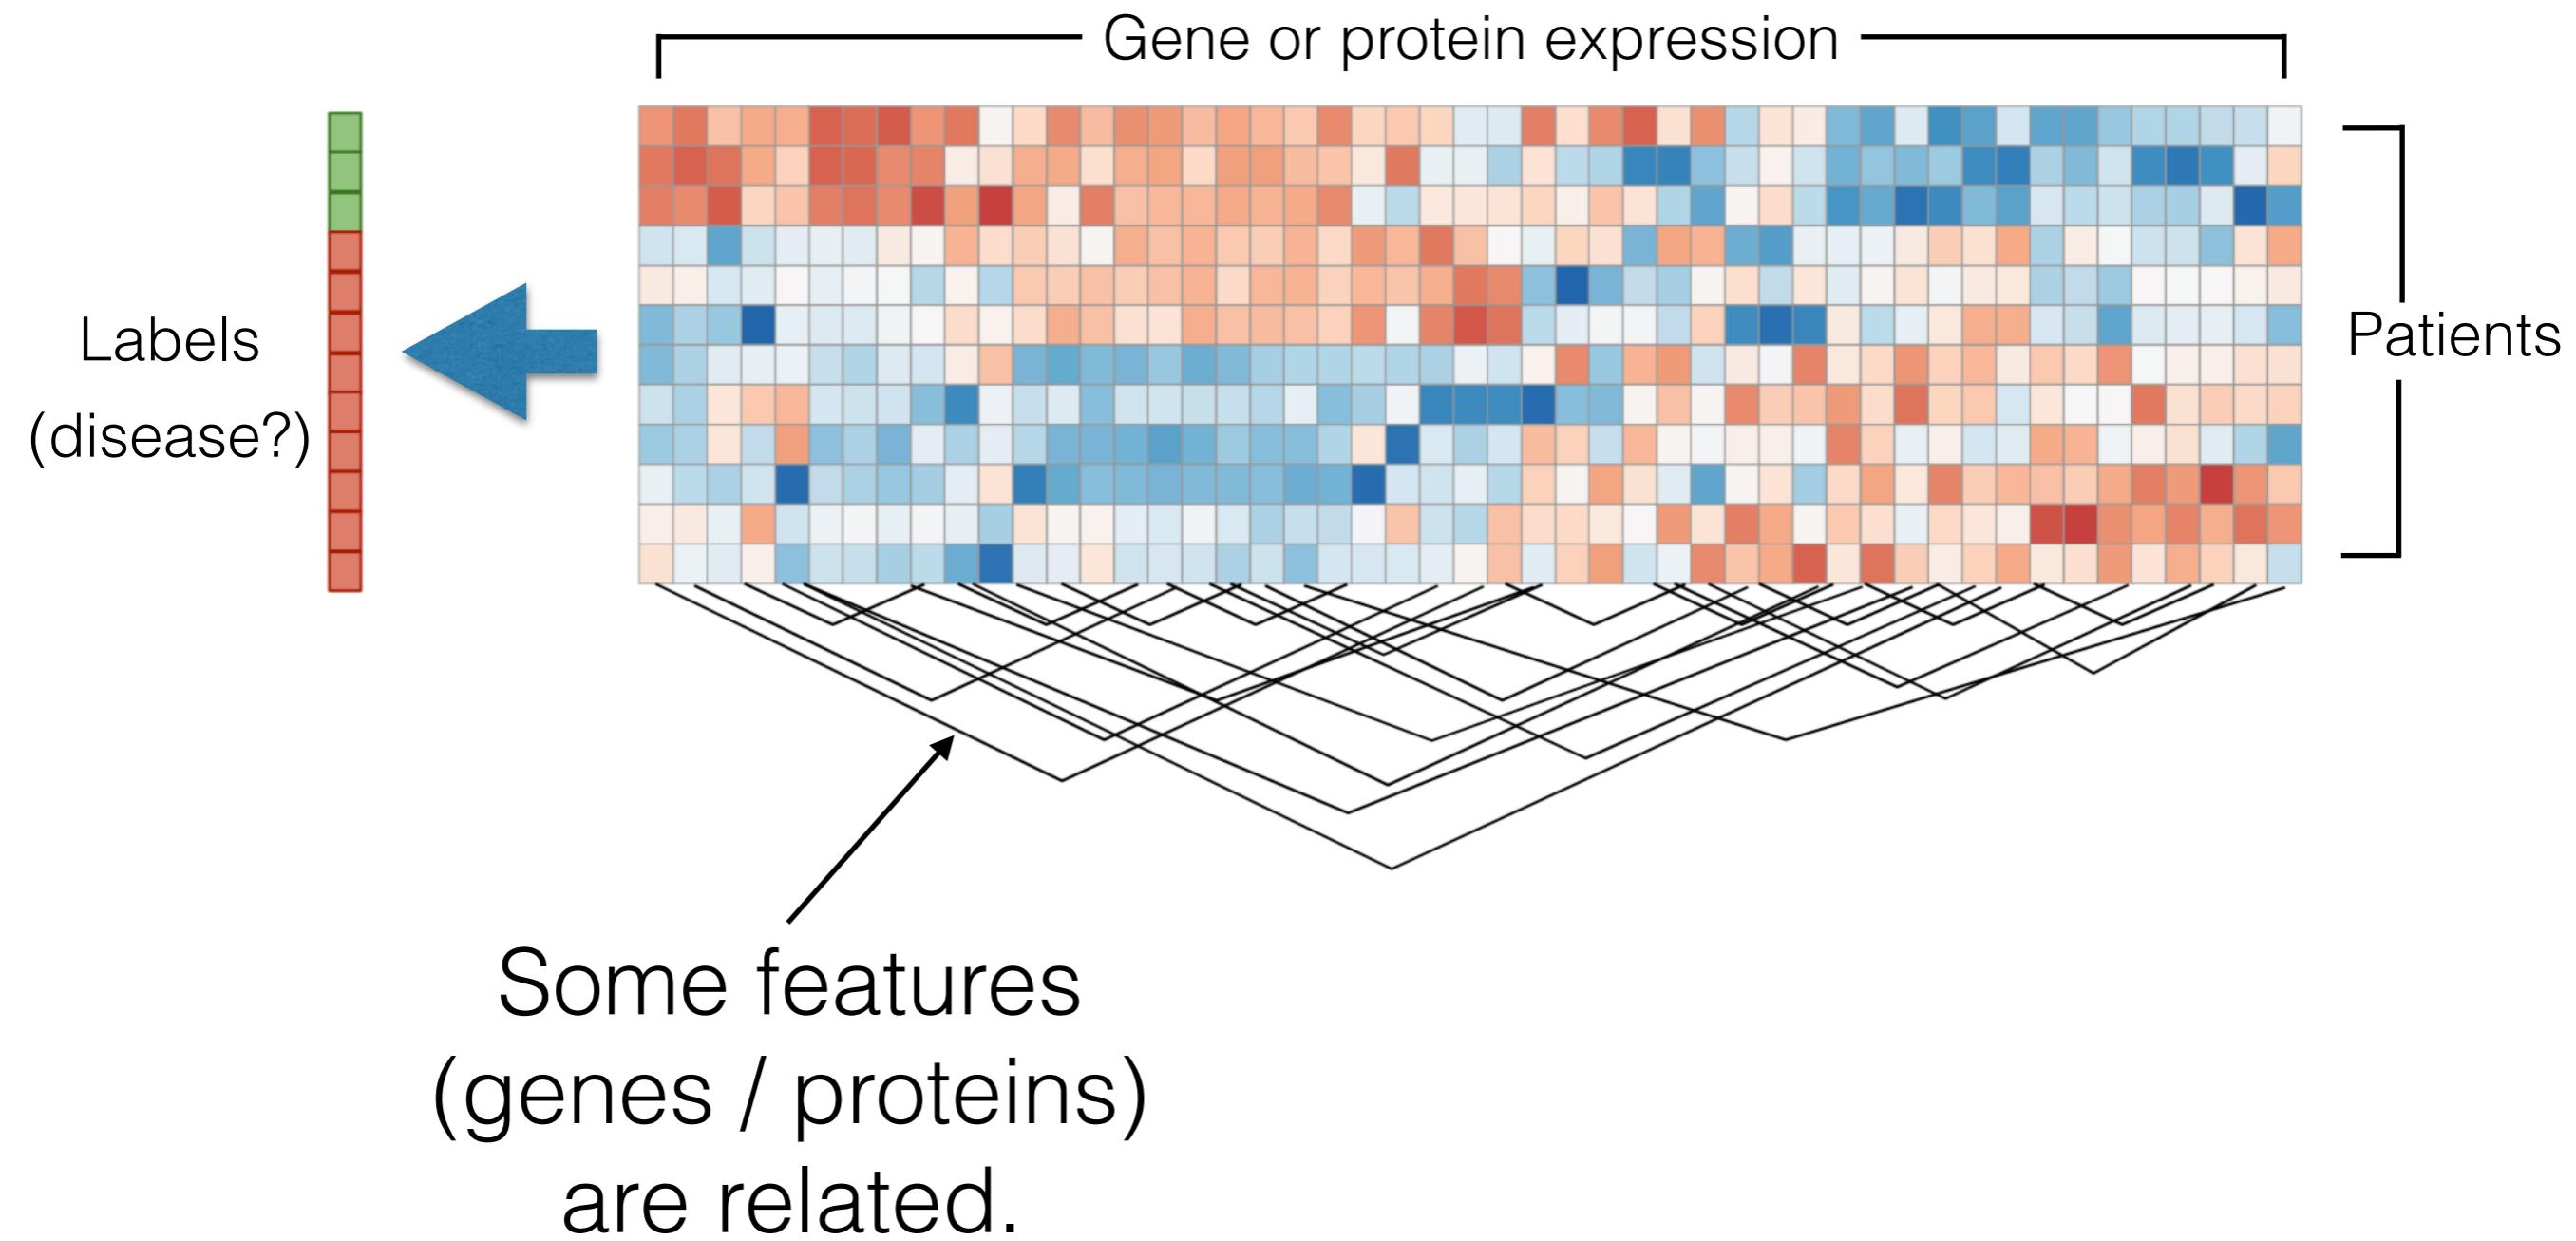
Example: Biological Network Data



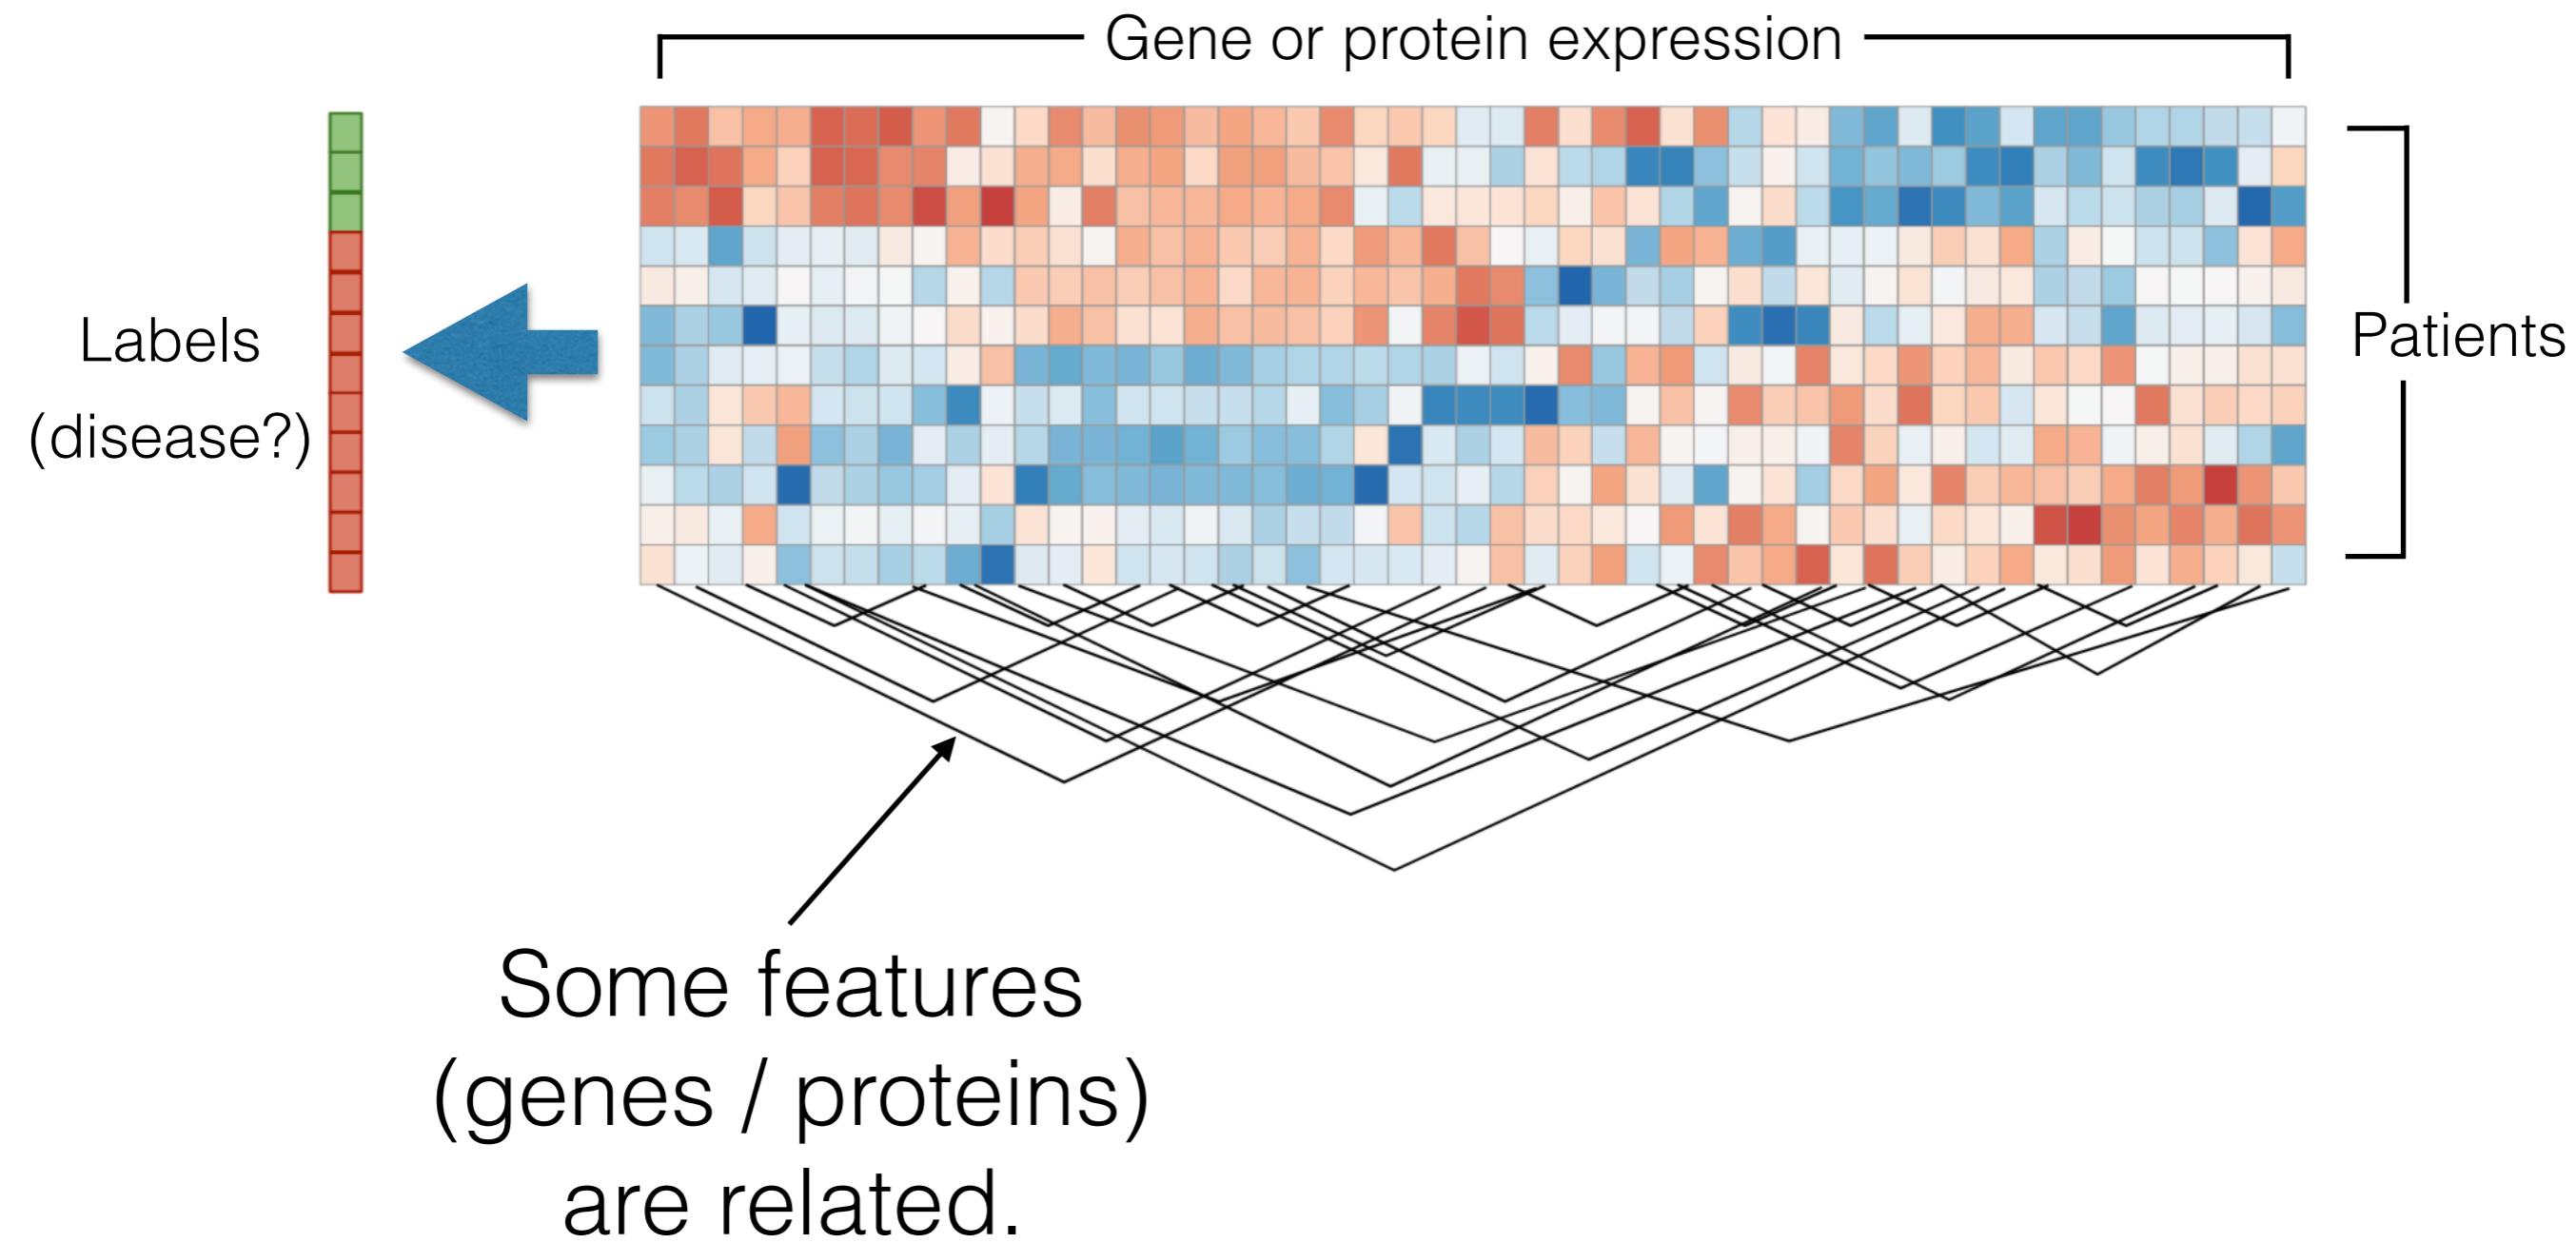
Example: Biological Network Data



Example: Biological Network Data



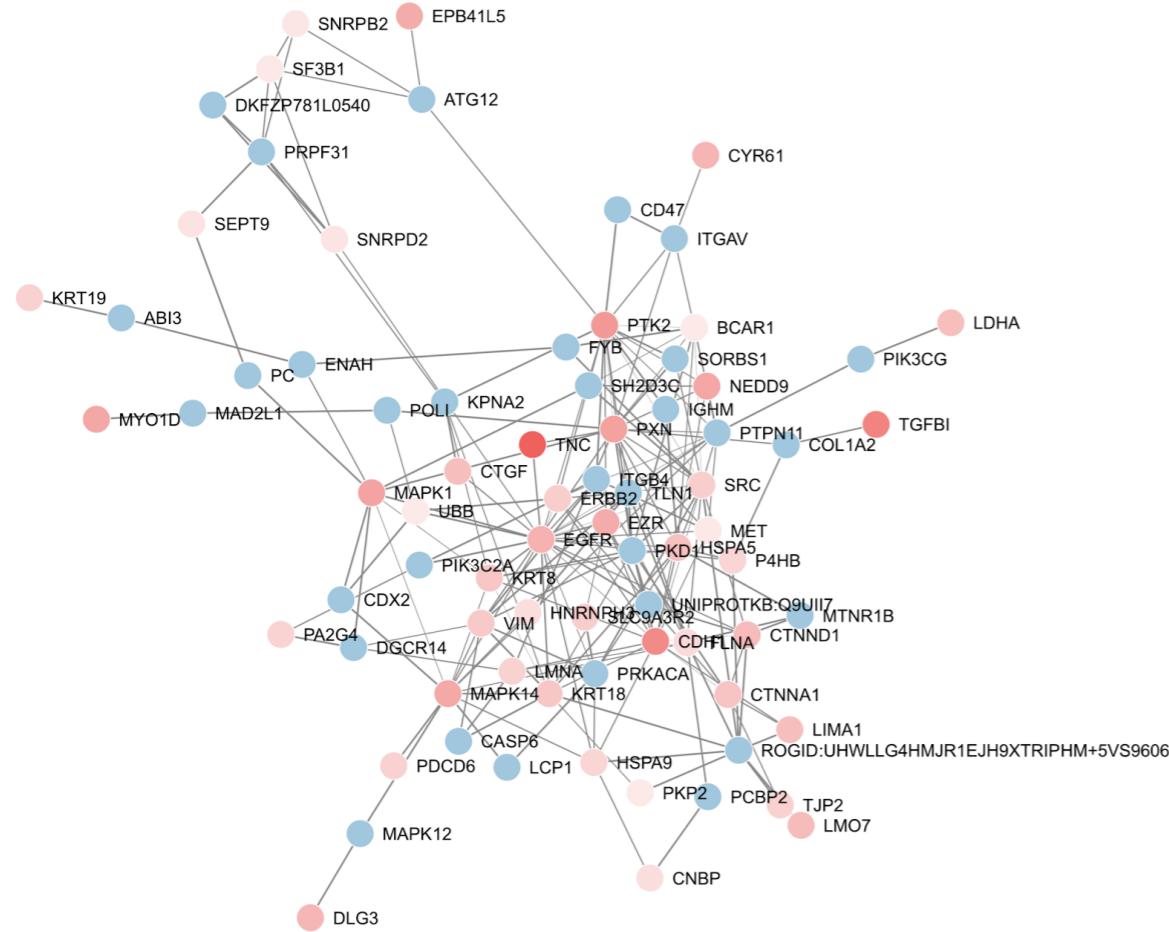
Example: Biological Network Data



Goal: find the most relevant genes / proteins.

Graph Structure

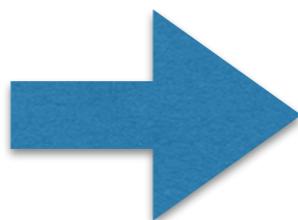
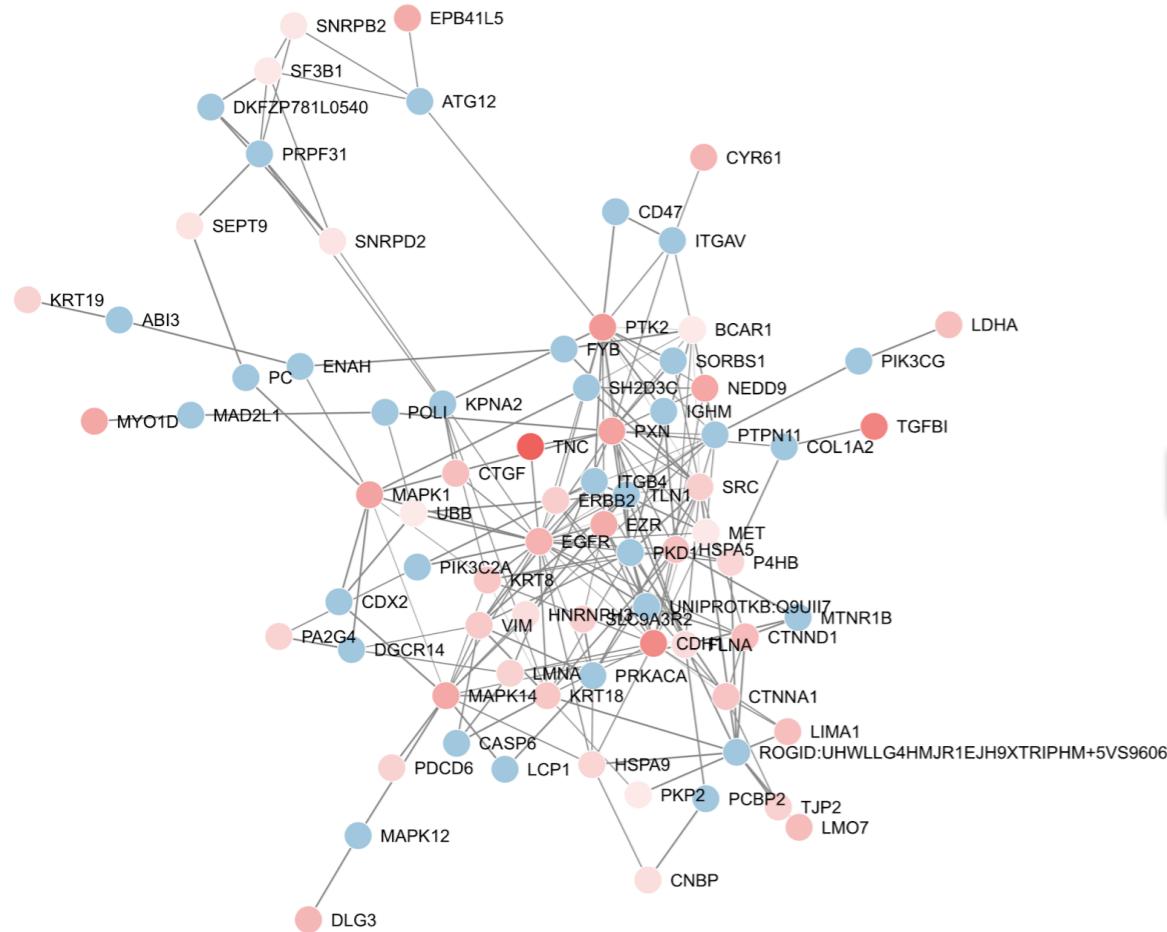
Input



Every node corresponds to a feature (gene / protein).

Graph Structure

Input

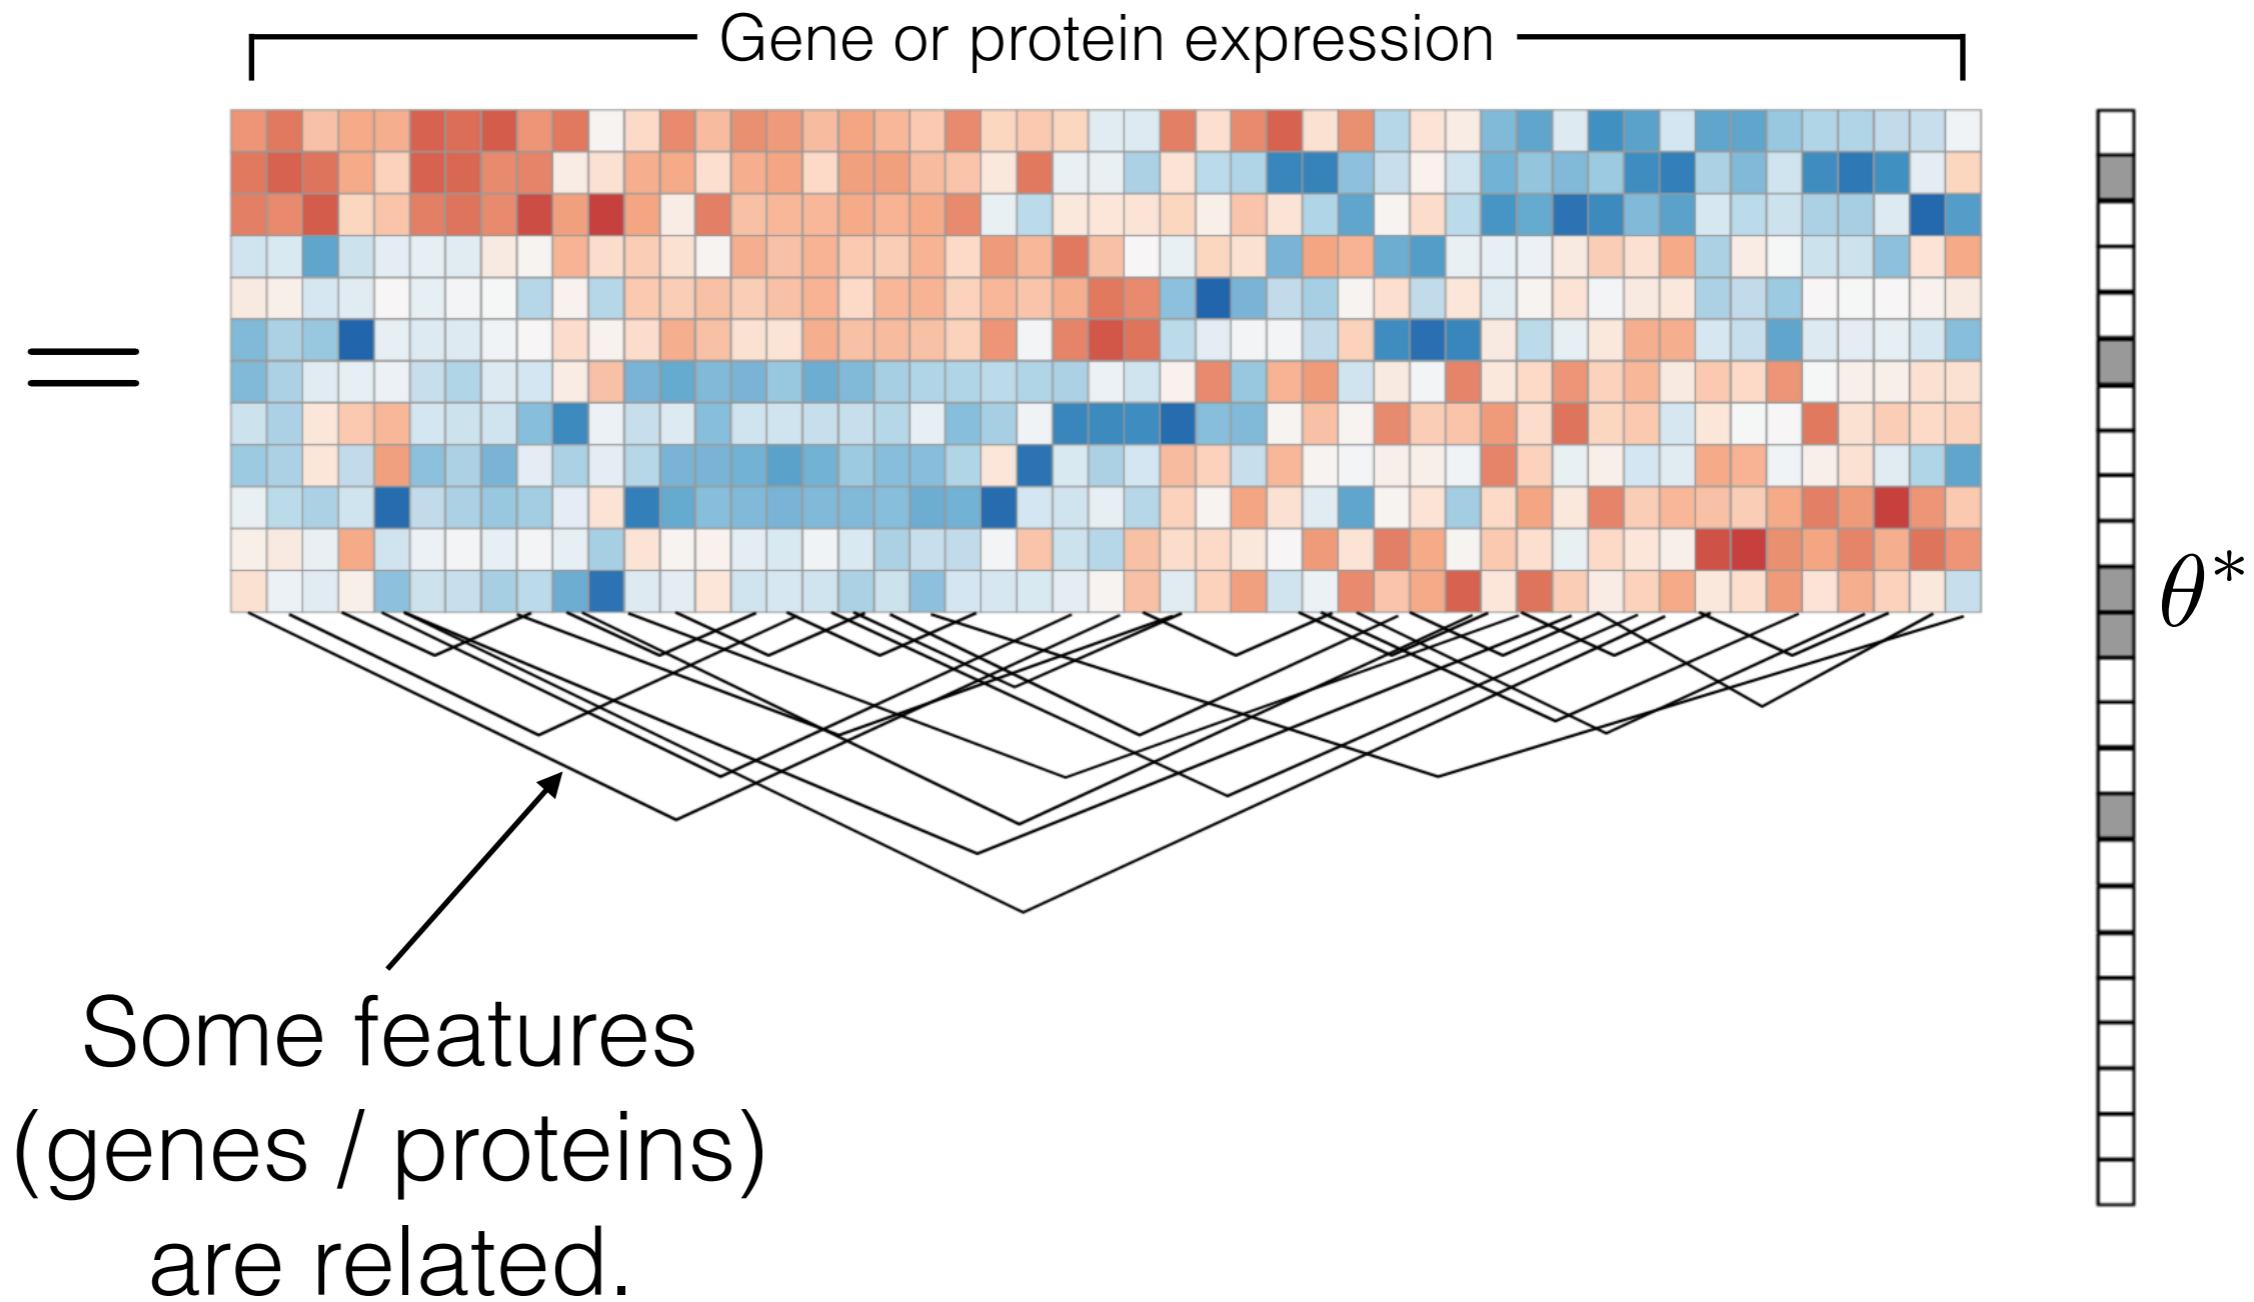


Output

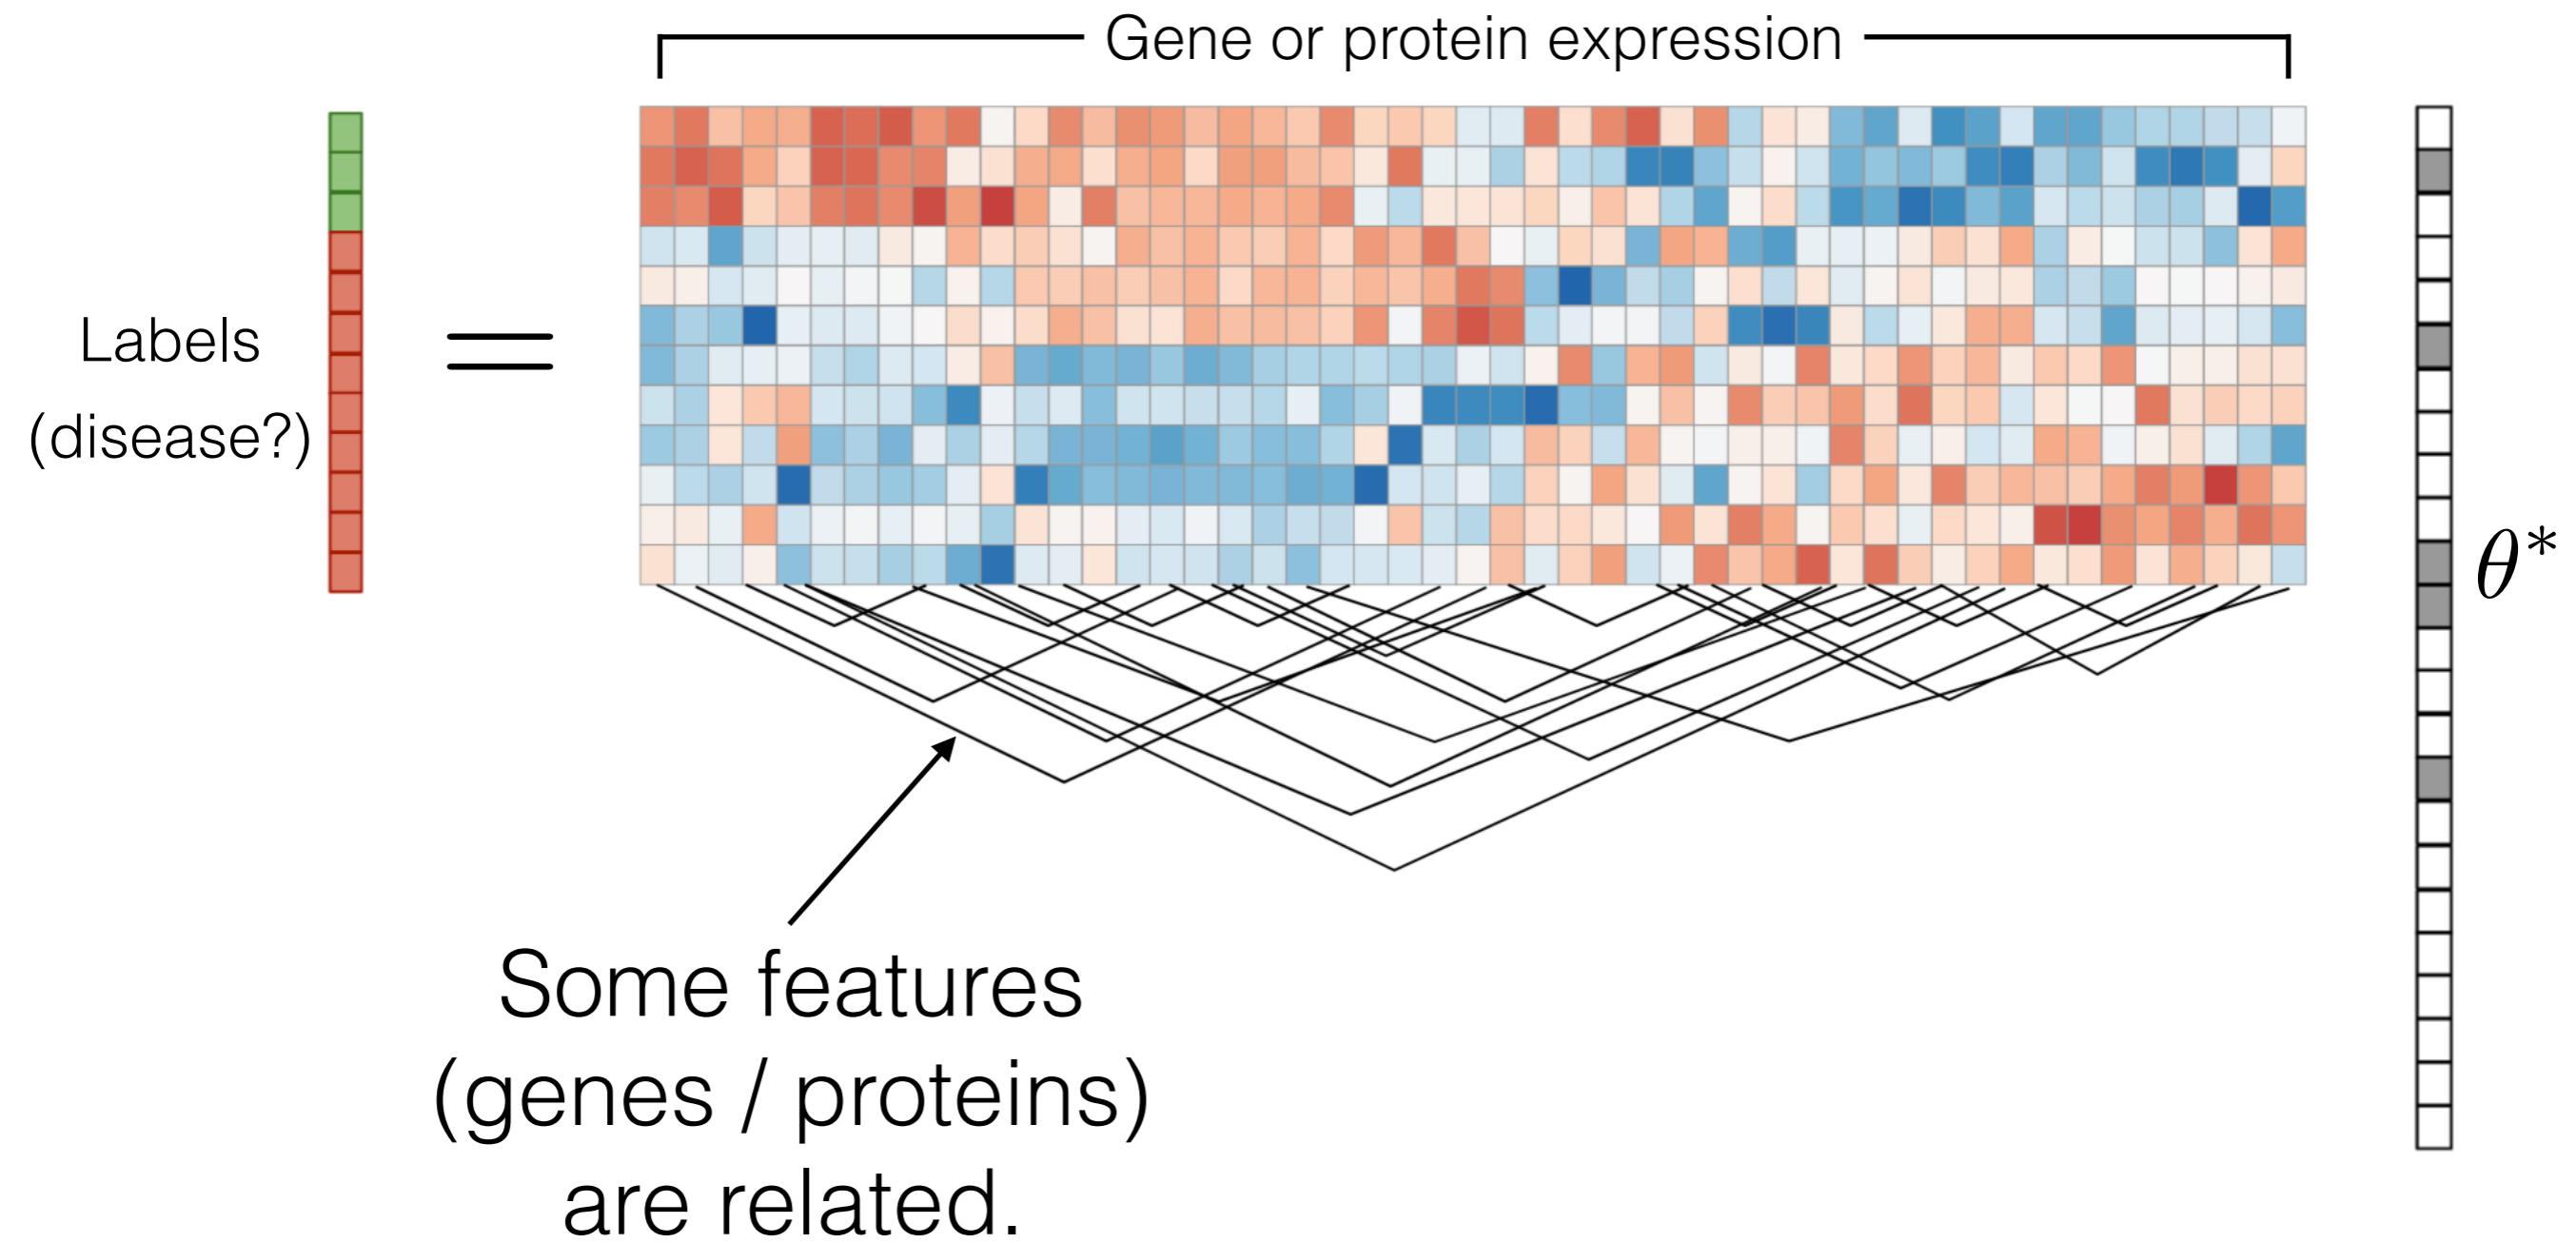
Every node corresponds to a feature (gene / protein).

Goal: Find one (or a few) connected subgraph with highly predictive features.

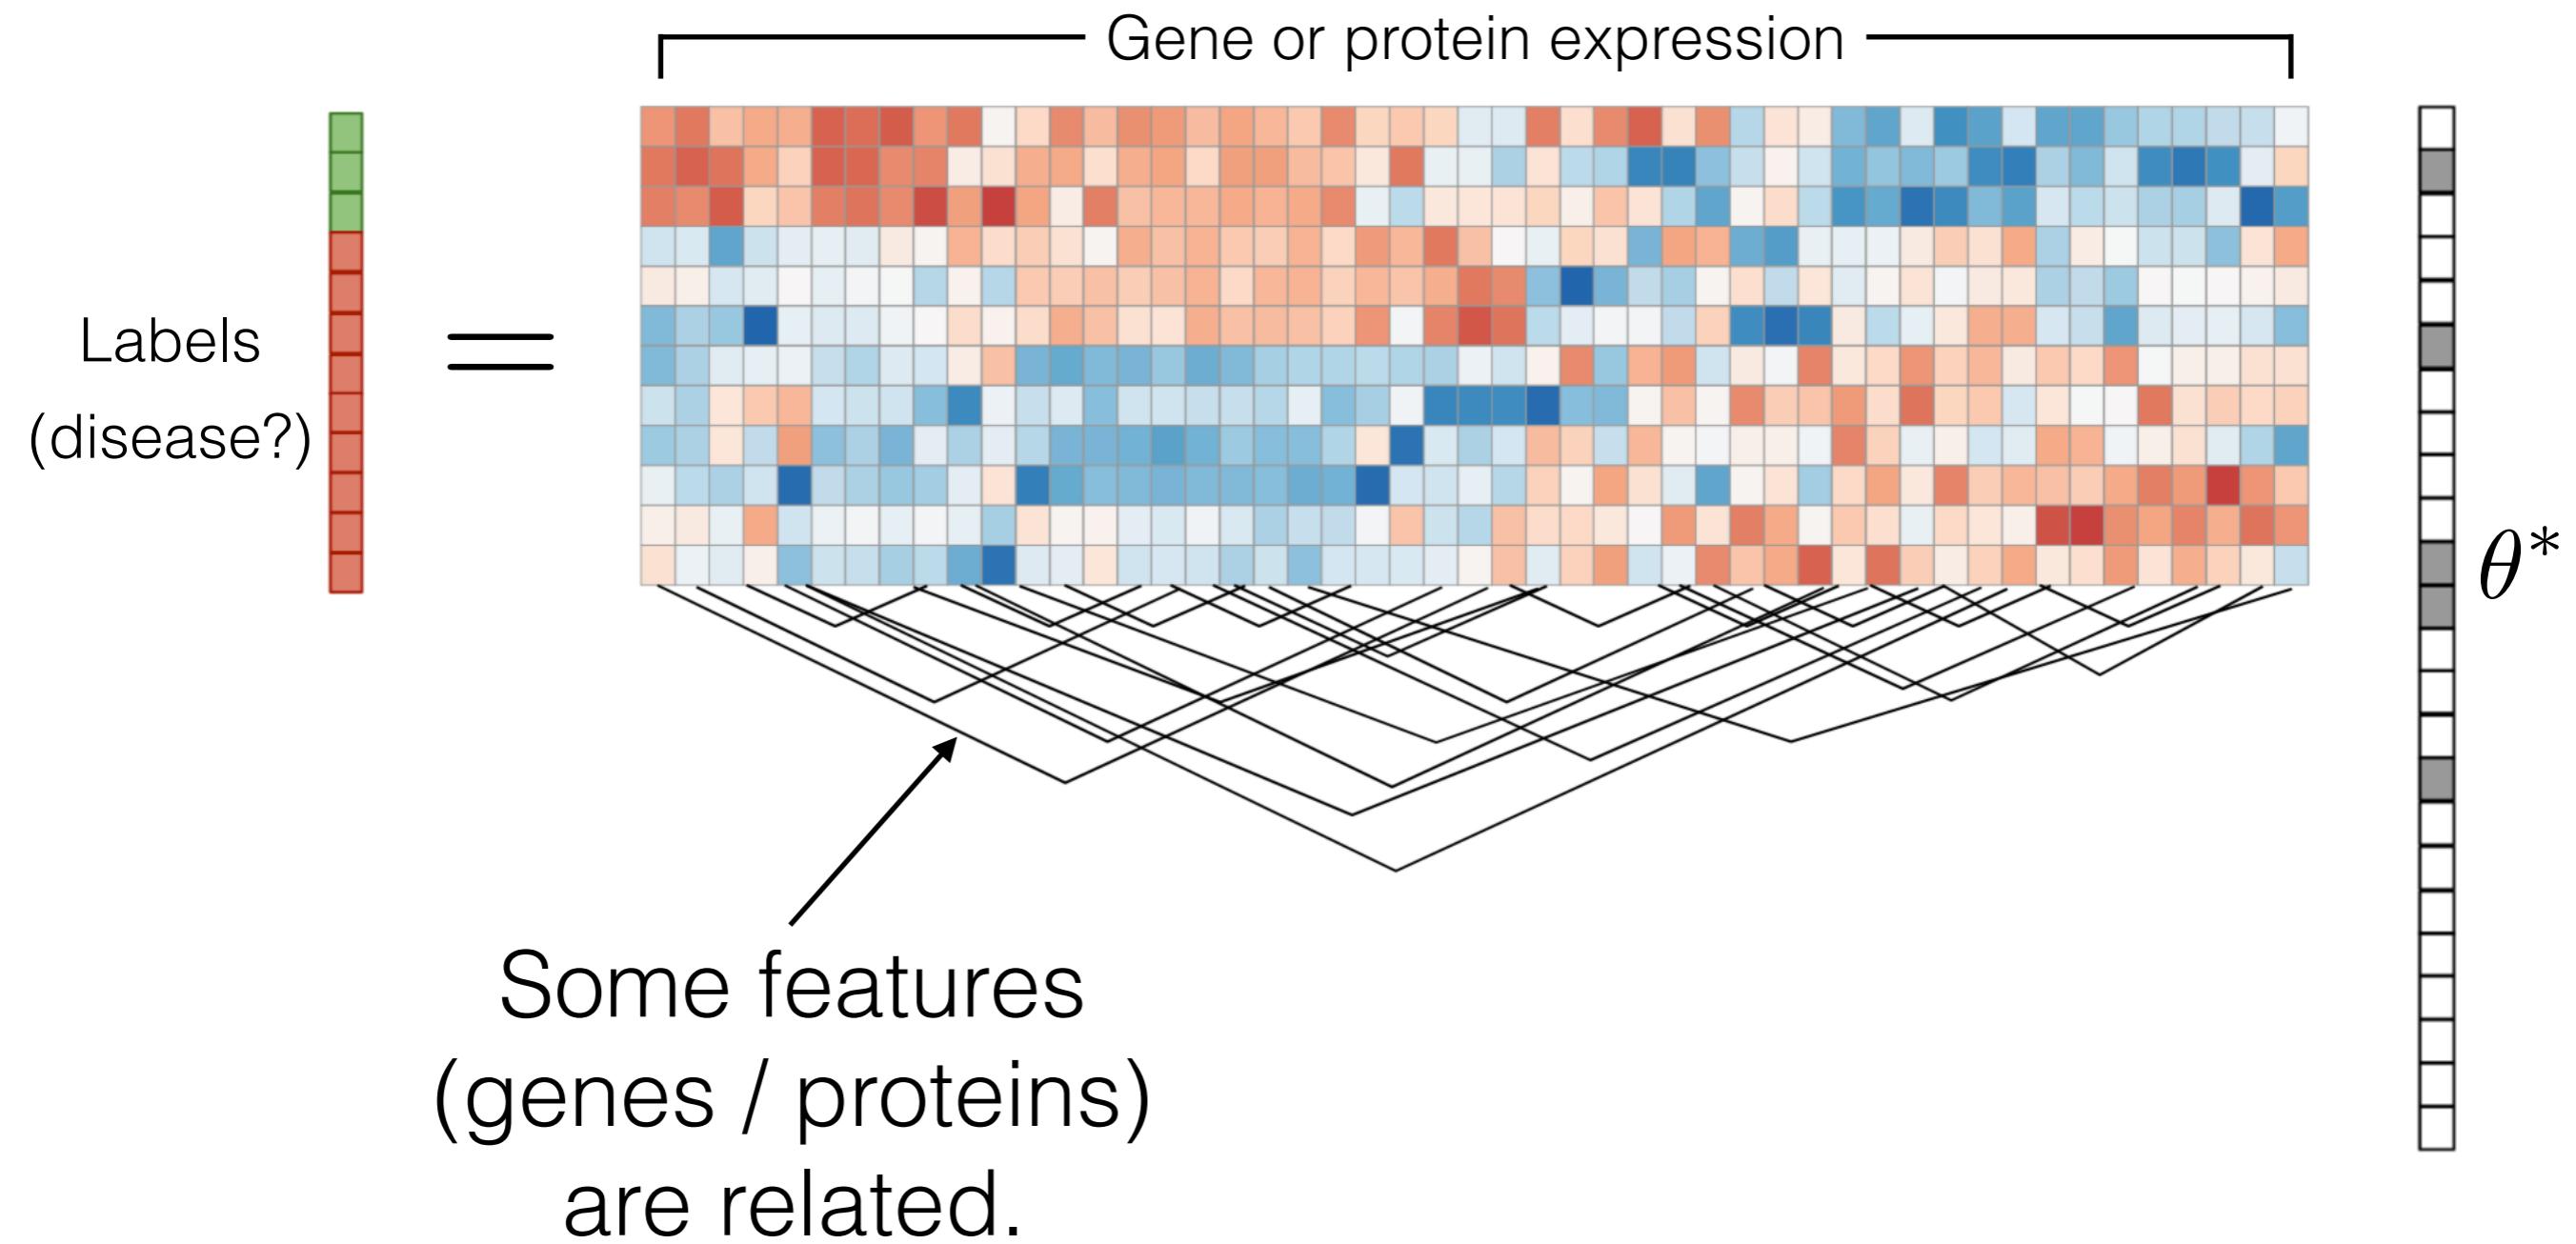
Linear Regression with Graph Constraint



Linear Regression with Graph Constraint



Linear Regression with Graph Constraint



Goal: find a **graph**-sparse set of parameters θ^* .

Solving Linear Regression with Constraints

	Sparse Linear Regression	Graph-Sparse Linear Regression
Algorithm	$\arg \min_{\ \theta\ _1 \leq \lambda} \ y - X\theta\ _2^2$	$\arg \min_{\theta \in \mathcal{C}} \ y - X\theta\ _2^2$
	The set \mathcal{C} is the set of graph-sparse vectors, i.e., vectors with a sparsity pattern corresponding to connected subgraphs.	

Solving Linear Regression with Constraints

	Sparse Linear Regression	Graph-Sparse Linear Regression
Algorithm	$\arg \min_{\ \theta\ _1 \leq \lambda} \ y - X\theta\ _2^2$	$\arg \min_{\theta \in \mathcal{C}} \ y - X\theta\ _2^2$
	The set \mathcal{C} is the set of graph-sparse vectors, i.e., vectors with a sparsity pattern corresponding to connected subgraphs.	

How can we solve this problem?

Projected Gradient Descent

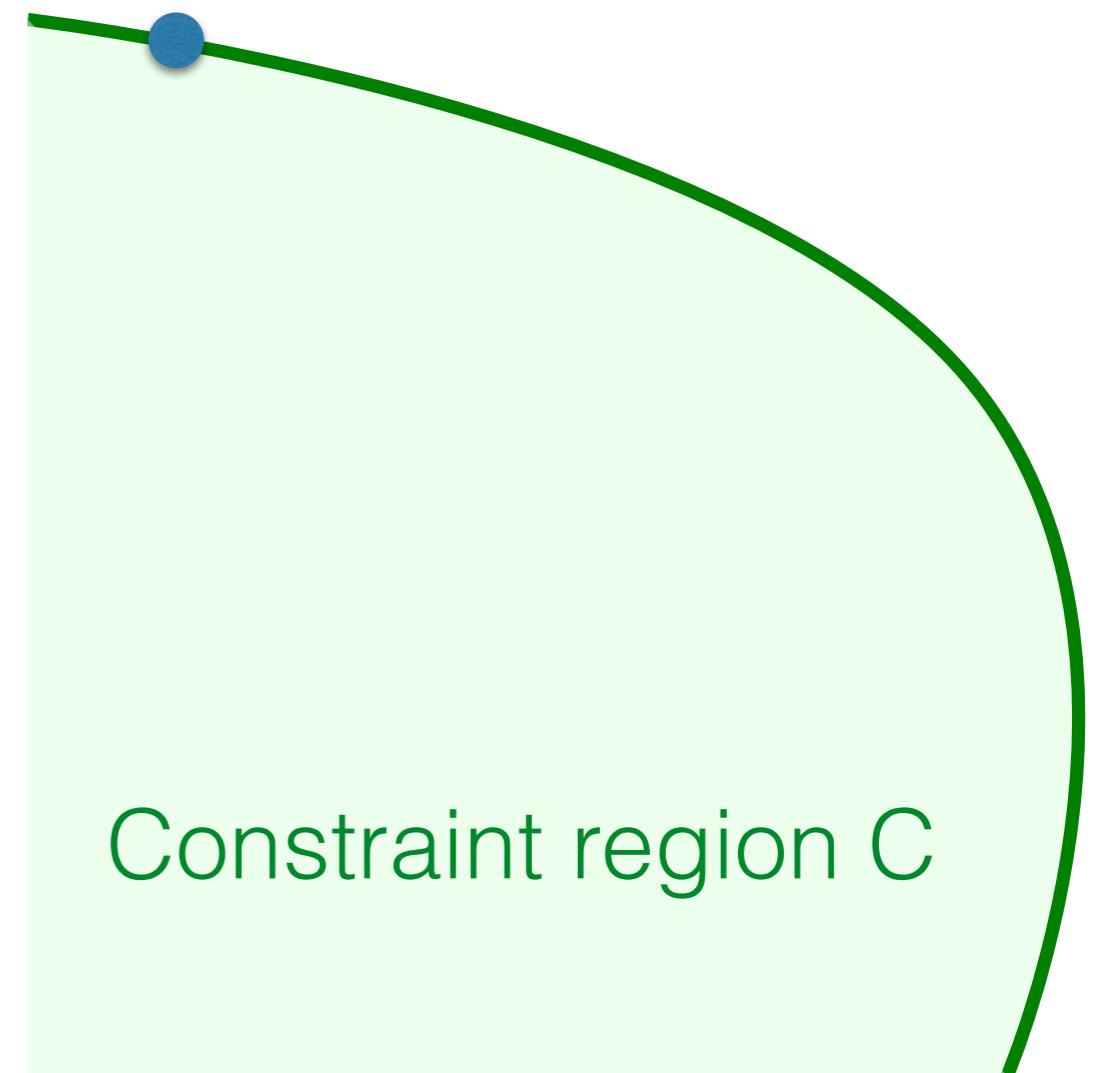
Repeat the following:

1. Compute gradient $\nabla f(\theta^i)$
2. Take a gradient step

$$\tilde{\theta}^{i+1} \leftarrow \theta^i + \eta \nabla f(\theta^i)$$

3. Project back to constraint set

$$\theta^{i+1} \leftarrow P_{\mathcal{C}}(\tilde{\theta}^{i+1})$$



Projected Gradient Descent

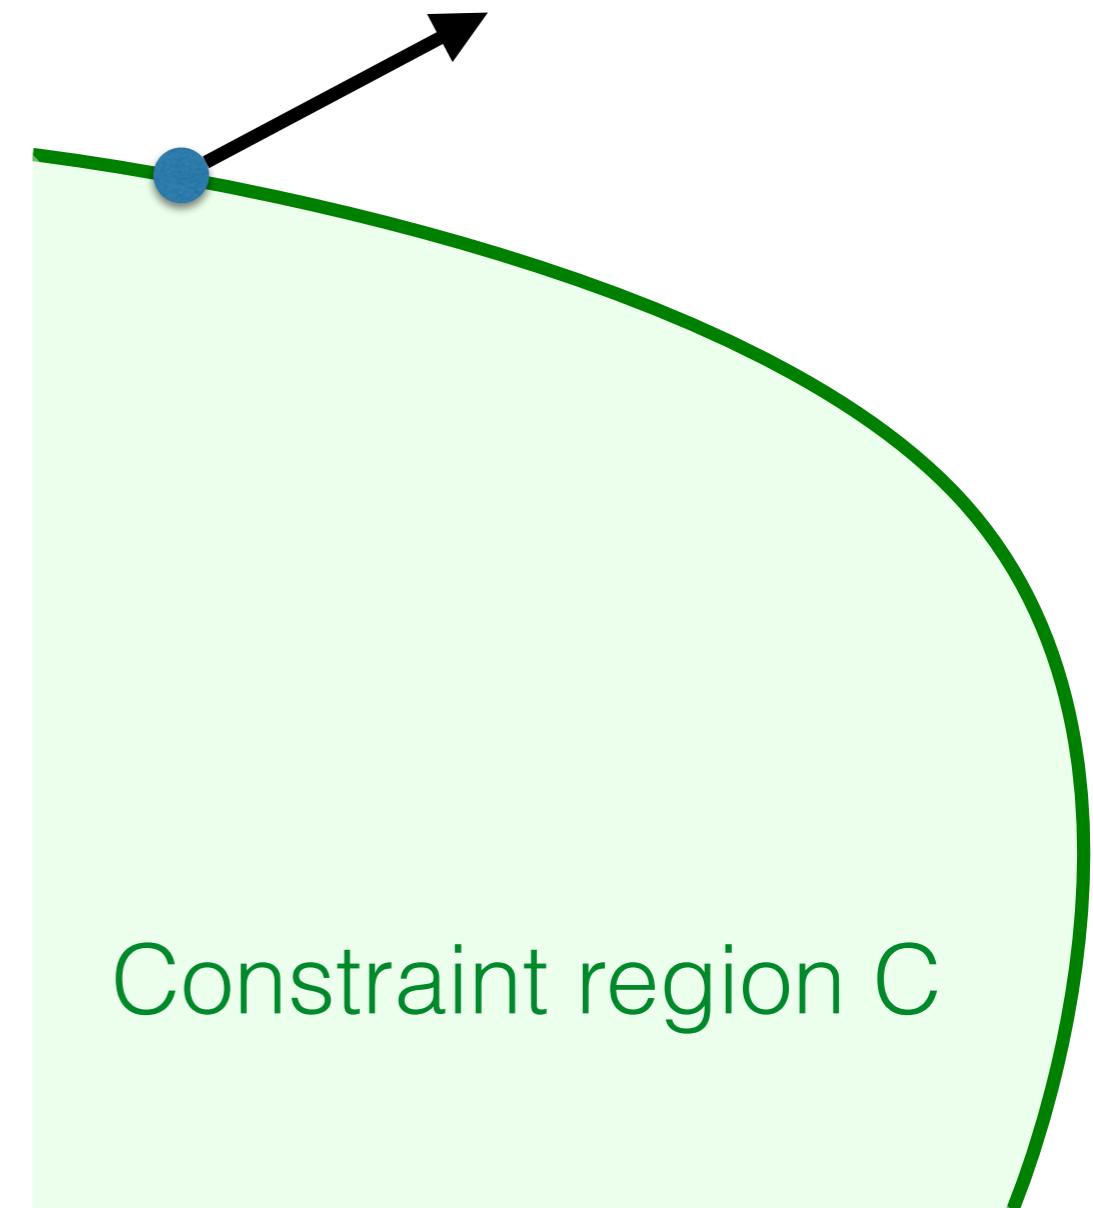
Repeat the following:

1. Compute gradient $\nabla f(\theta^i)$
2. Take a gradient step

$$\tilde{\theta}^{i+1} \leftarrow \theta^i + \eta \nabla f(\theta^i)$$

3. Project back to constraint set

$$\theta^{i+1} \leftarrow P_{\mathcal{C}}(\tilde{\theta}^{i+1})$$



Projected Gradient Descent

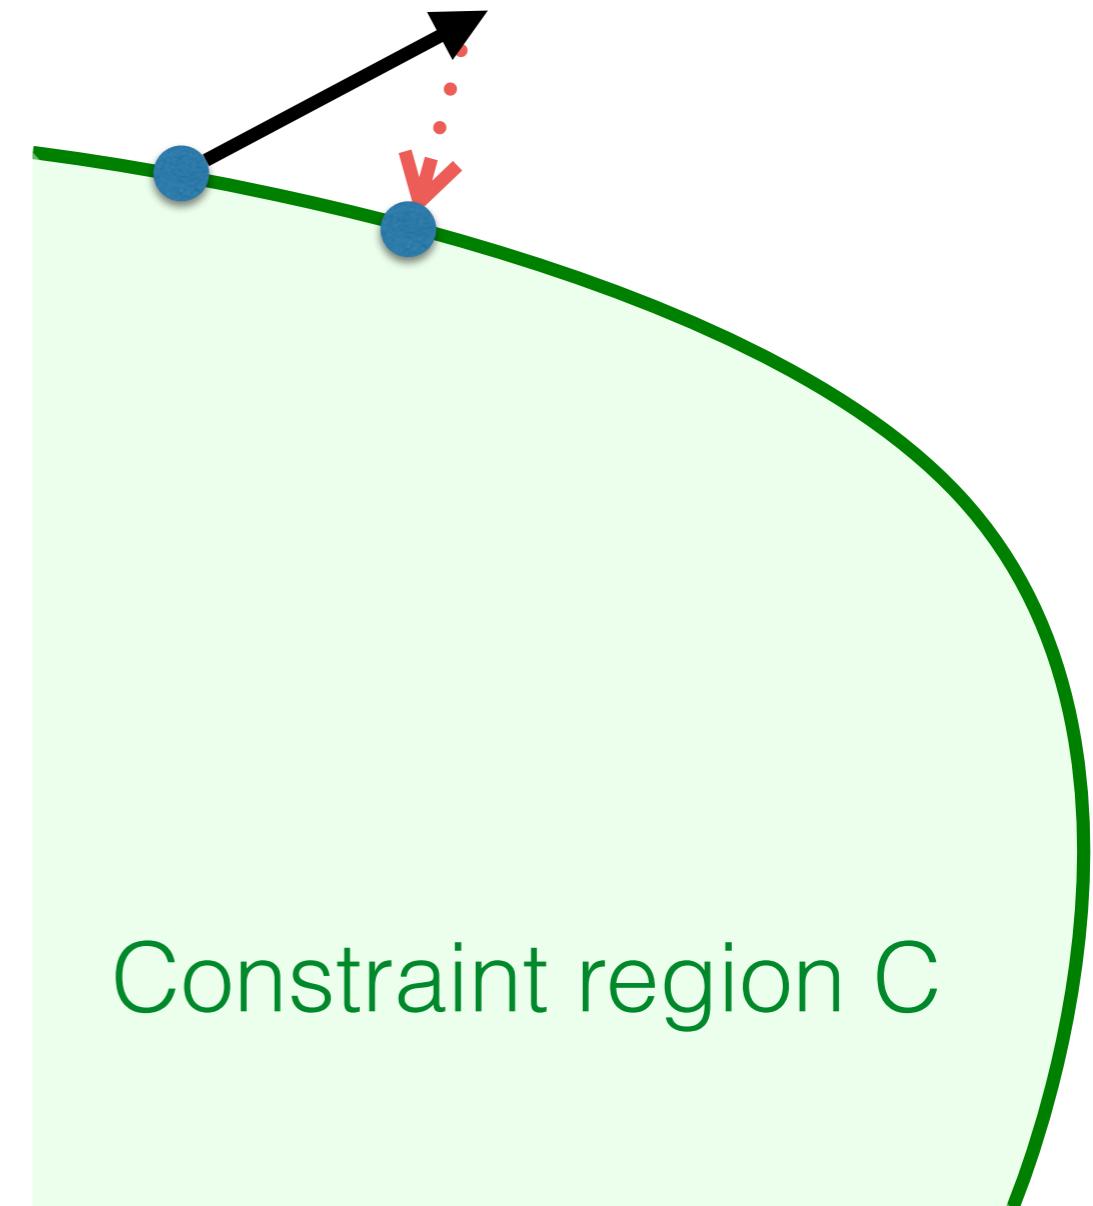
Repeat the following:

1. Compute gradient $\nabla f(\theta^i)$
2. Take a gradient step

$$\tilde{\theta}^{i+1} \leftarrow \theta^i + \eta \nabla f(\theta^i)$$

3. Project back to constraint set

$$\theta^{i+1} \leftarrow P_{\mathcal{C}}(\tilde{\theta}^{i+1})$$



Projected Gradient Descent

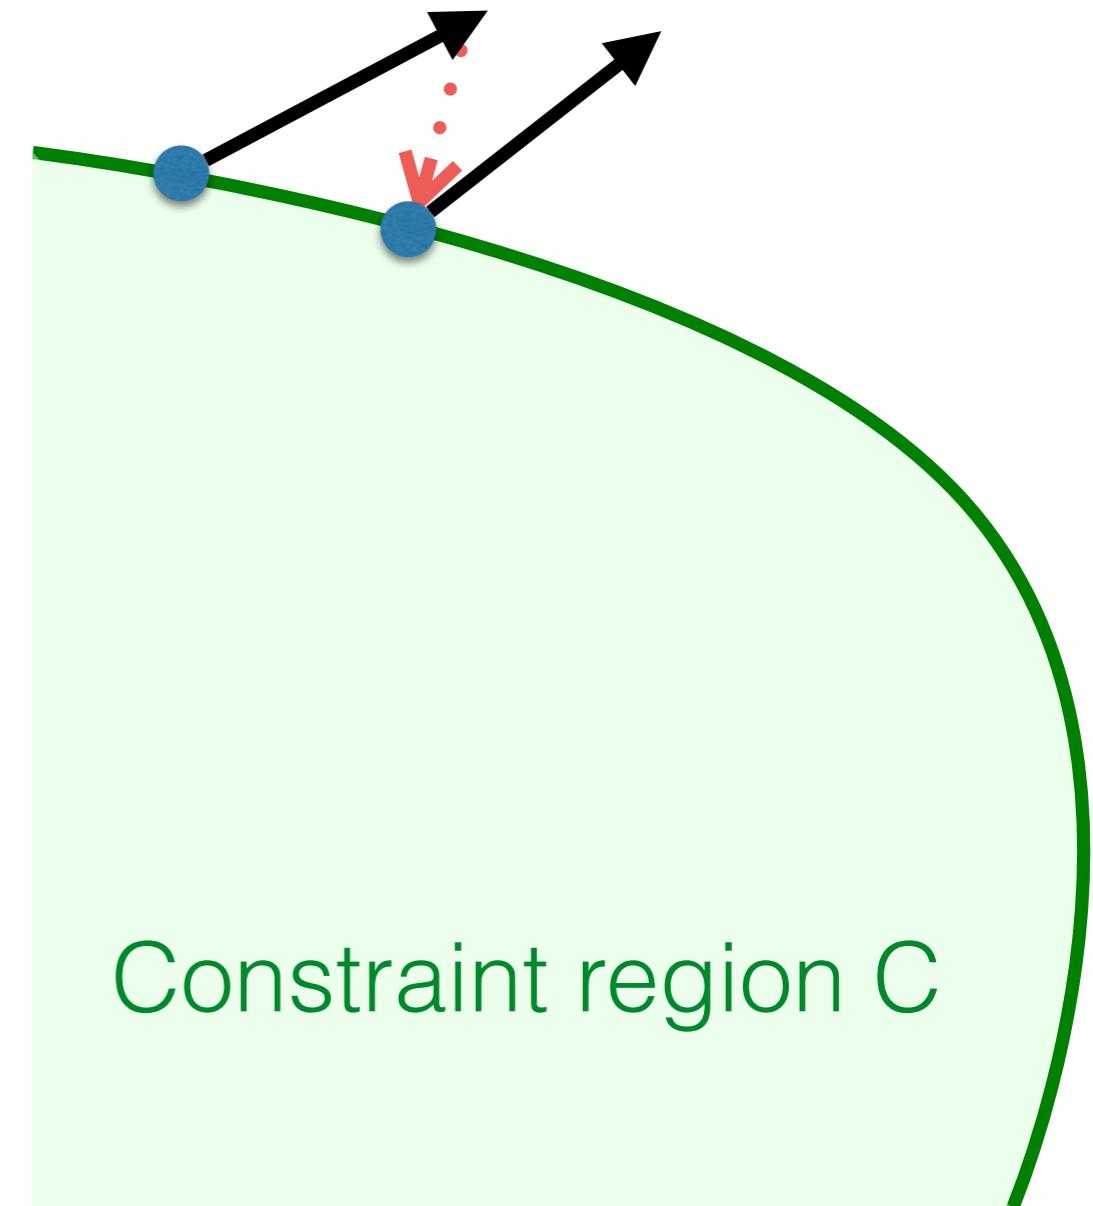
Repeat the following:

1. Compute gradient $\nabla f(\theta^i)$
2. Take a gradient step

$$\tilde{\theta}^{i+1} \leftarrow \theta^i + \eta \nabla f(\theta^i)$$

3. Project back to constraint set

$$\theta^{i+1} \leftarrow P_{\mathcal{C}}(\tilde{\theta}^{i+1})$$



Projected Gradient Descent

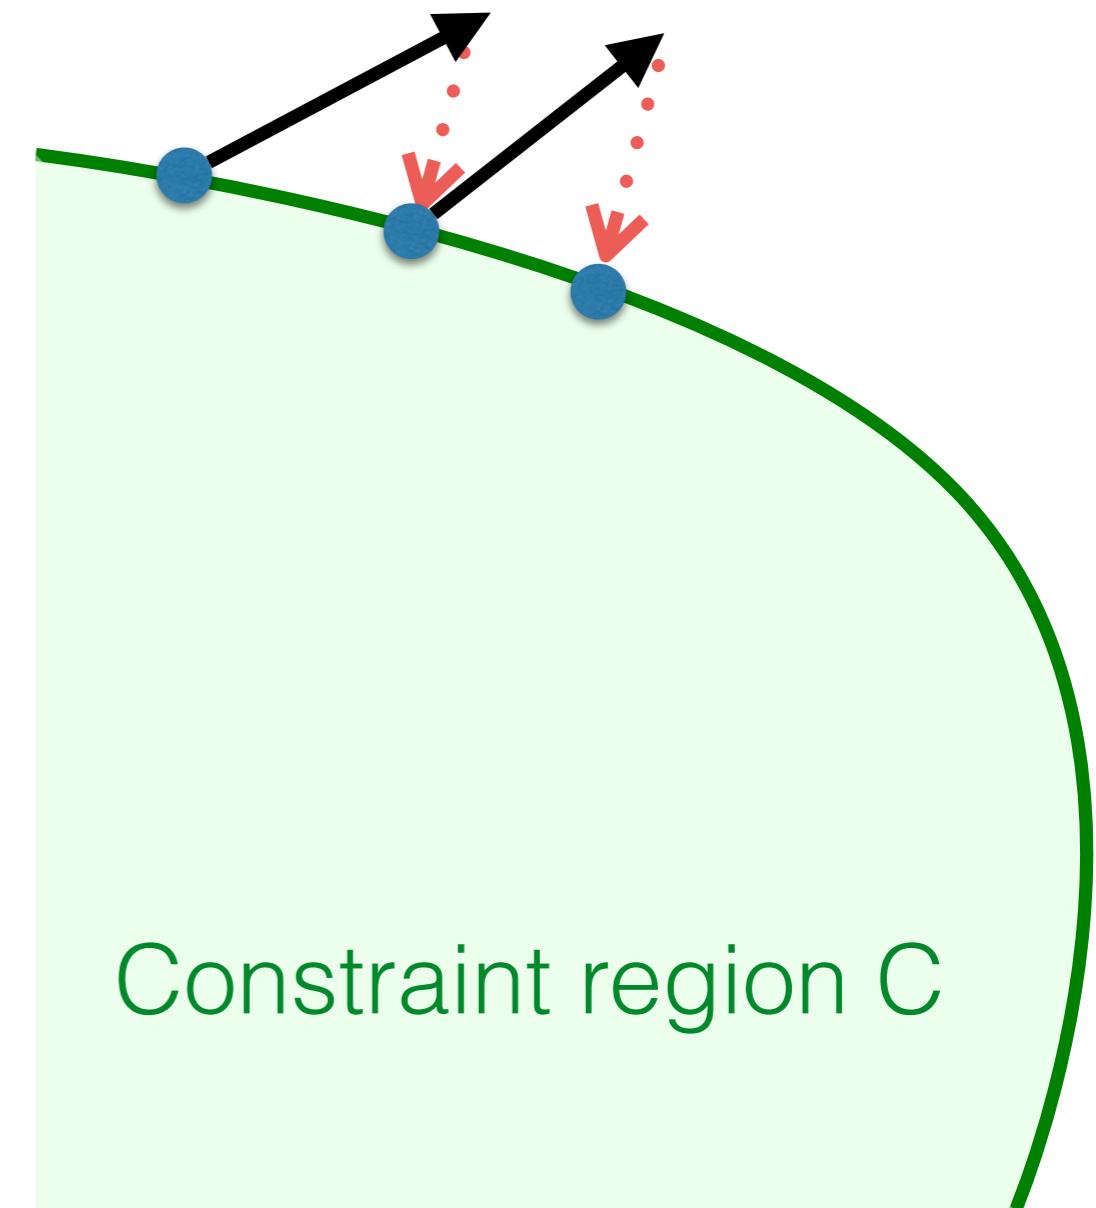
Repeat the following:

1. Compute gradient $\nabla f(\theta^i)$
2. Take a gradient step

$$\tilde{\theta}^{i+1} \leftarrow \theta^i + \eta \nabla f(\theta^i)$$

3. Project back to constraint set

$$\theta^{i+1} \leftarrow P_{\mathcal{C}}(\tilde{\theta}^{i+1})$$



Projected Gradient Descent

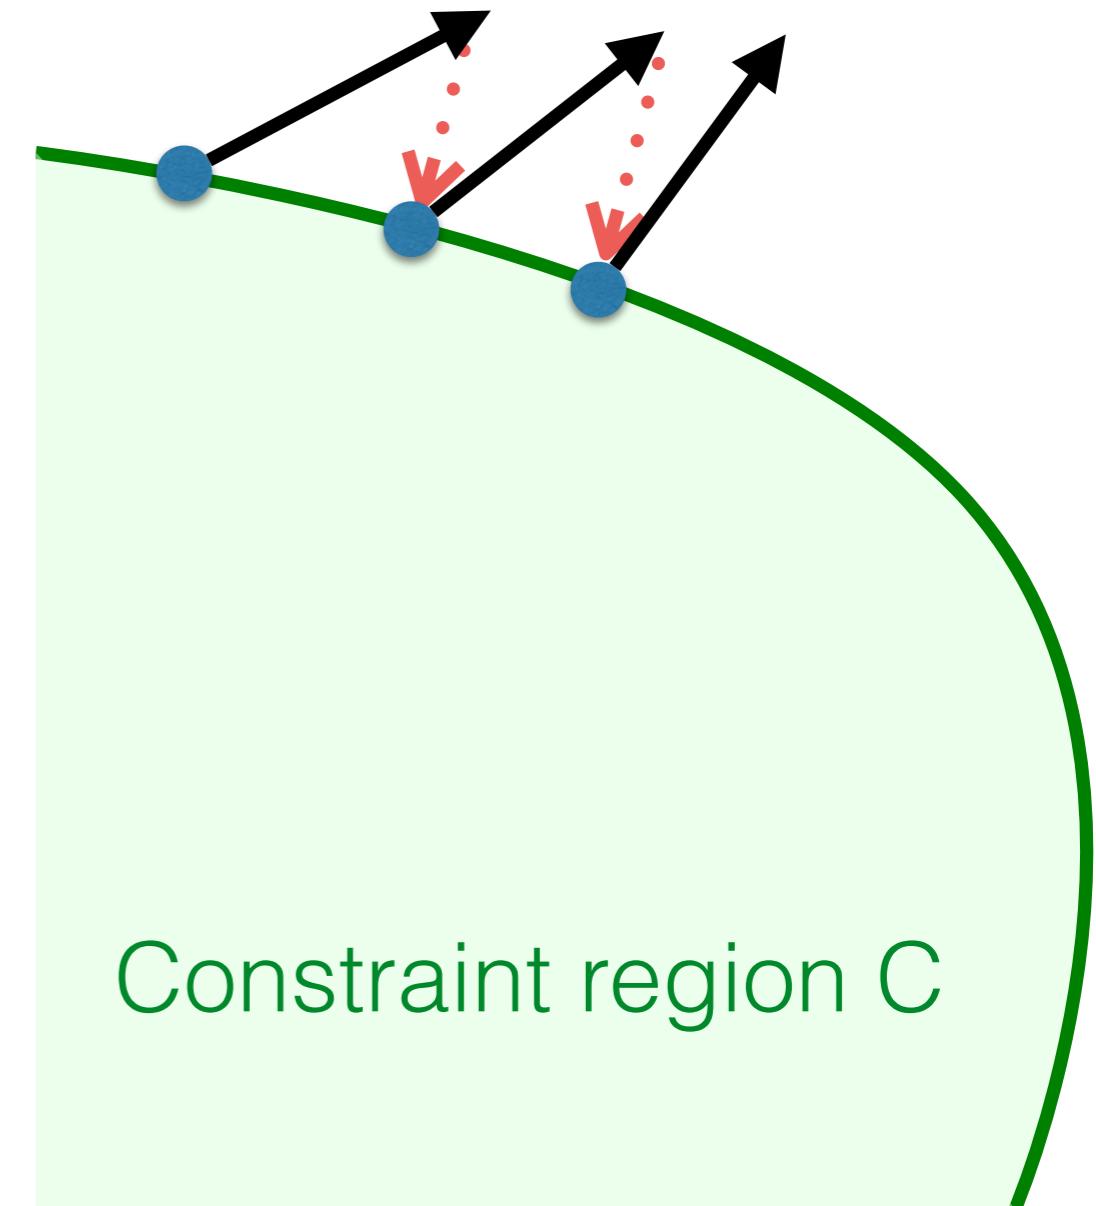
Repeat the following:

1. Compute gradient $\nabla f(\theta^i)$
2. Take a gradient step

$$\tilde{\theta}^{i+1} \leftarrow \theta^i + \eta \nabla f(\theta^i)$$

3. Project back to constraint set

$$\theta^{i+1} \leftarrow P_{\mathcal{C}}(\tilde{\theta}^{i+1})$$



Projected Gradient Descent

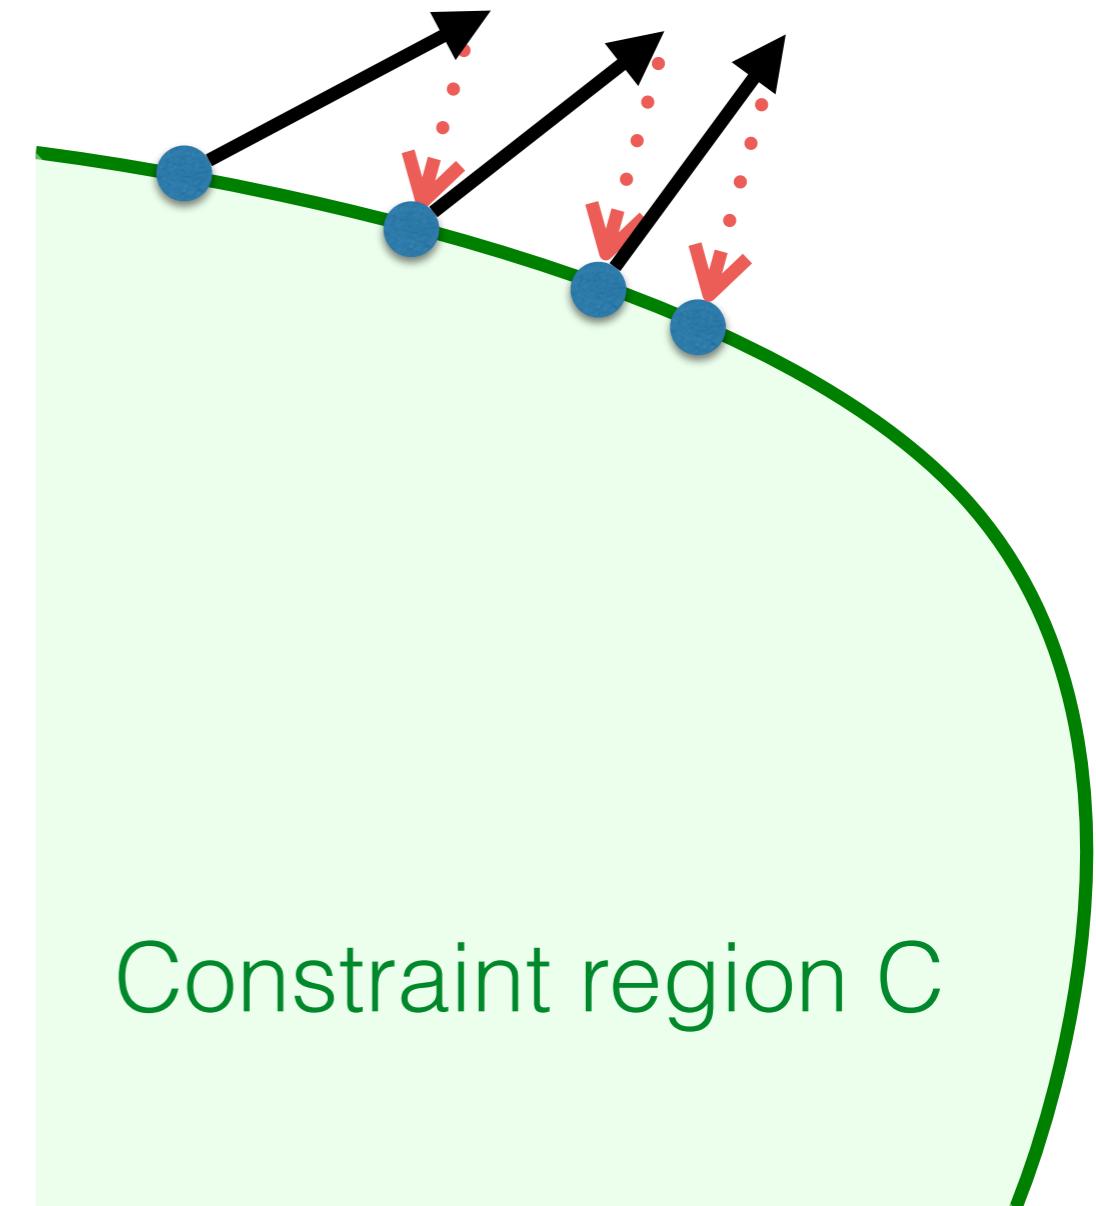
Repeat the following:

1. Compute gradient $\nabla f(\theta^i)$
2. Take a gradient step

$$\tilde{\theta}^{i+1} \leftarrow \theta^i + \eta \nabla f(\theta^i)$$

3. Project back to constraint set

$$\theta^{i+1} \leftarrow P_{\mathcal{C}}(\tilde{\theta}^{i+1})$$



Projected Gradient Descent

Repeat the following:

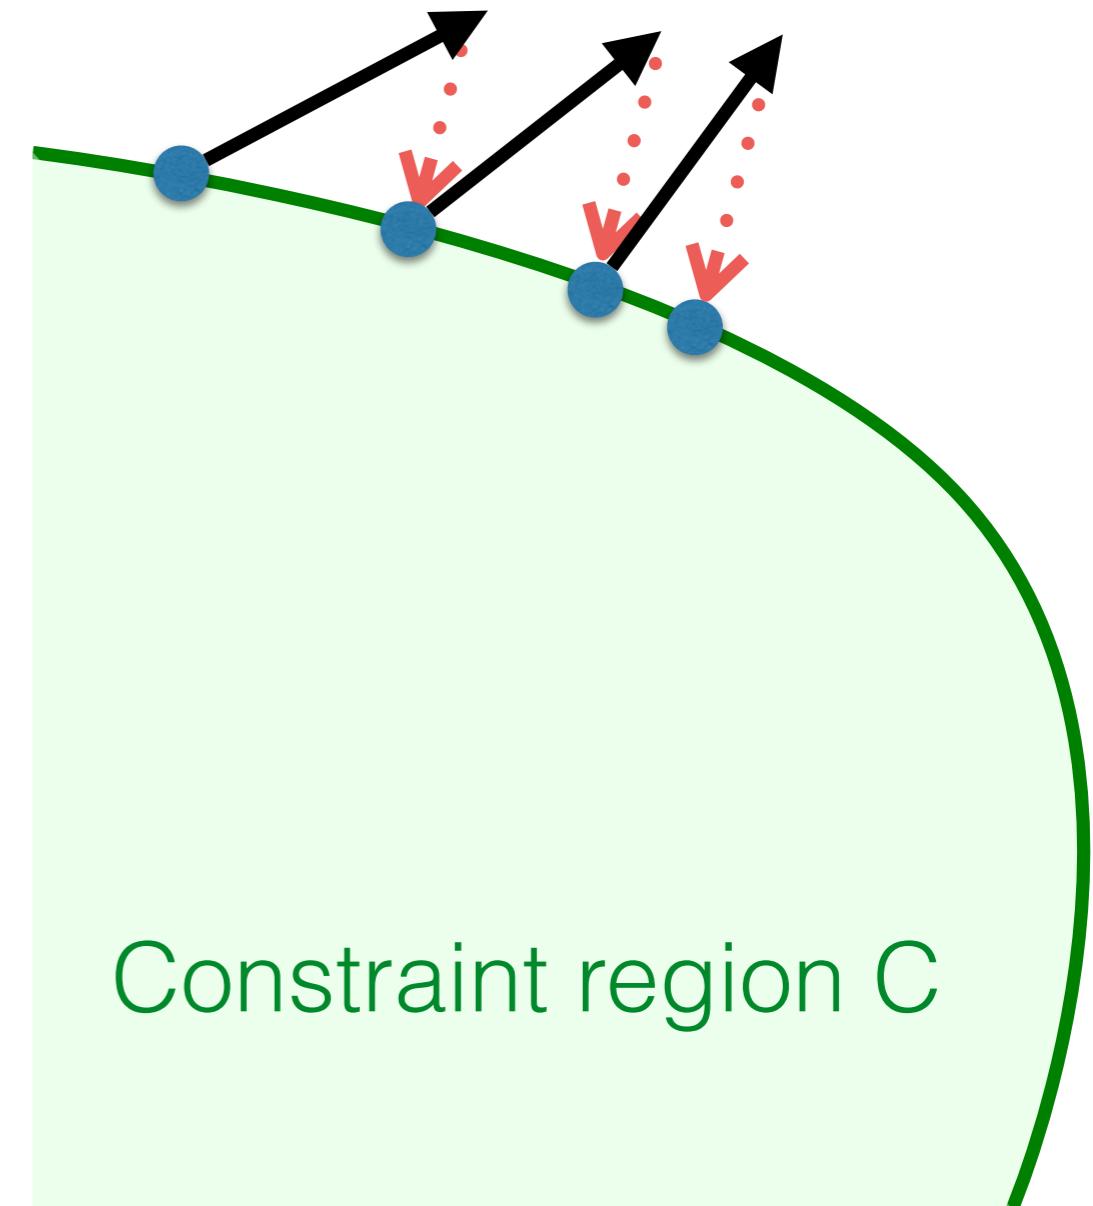
1. Compute gradient $\nabla f(\theta^i)$
2. Take a gradient step

$$\tilde{\theta}^{i+1} \leftarrow \theta^i + \eta \nabla f(\theta^i)$$

3. Project back to constraint set

$$\theta^{i+1} \leftarrow P_C(\tilde{\theta}^{i+1})$$

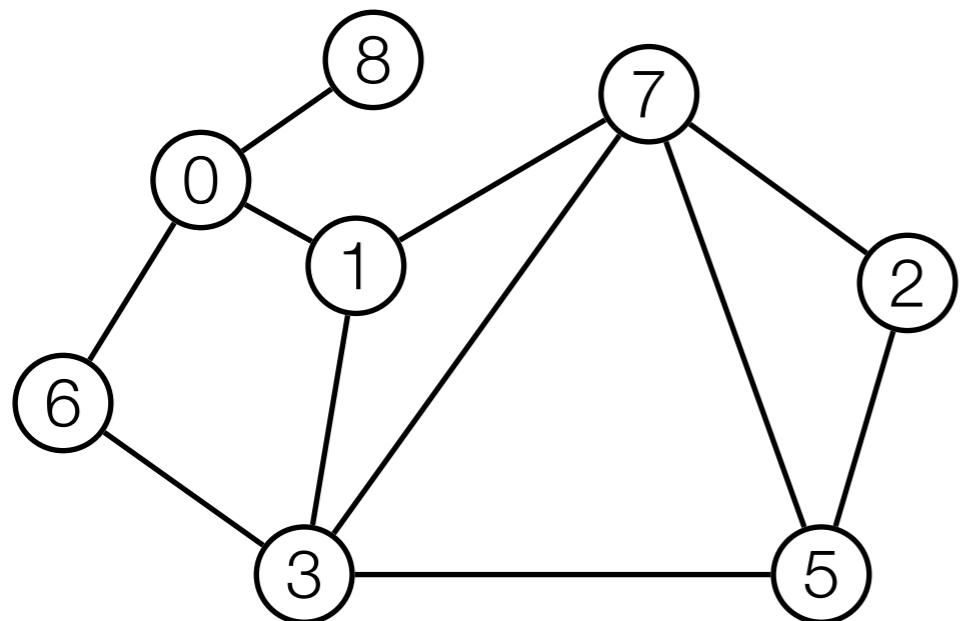
Computational bottleneck



Graph Sparsity Projection

Given a node-weighted graph, find a set of g subtrees with total size s so that the sum of node weights is maximized.

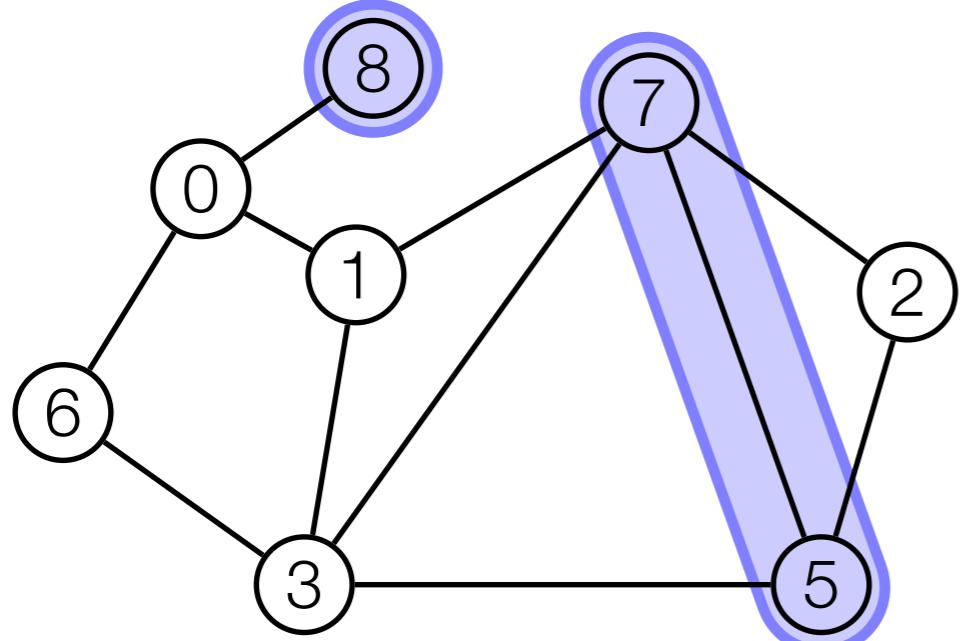
Example: $g = 2, s = 3$



Graph Sparsity Projection

Given a node-weighted graph, find a set of g subtrees with total size s so that the sum of node weights is maximized.

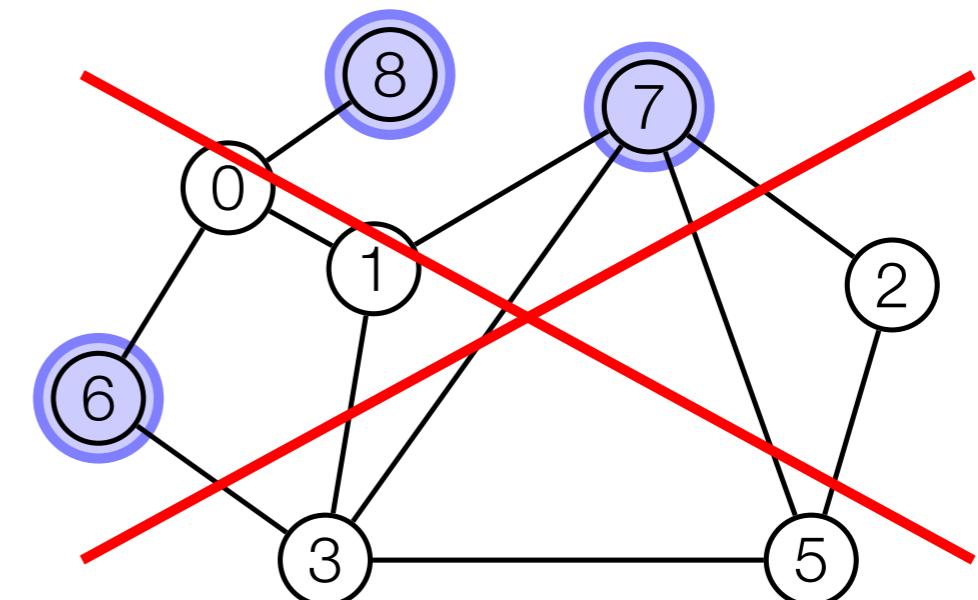
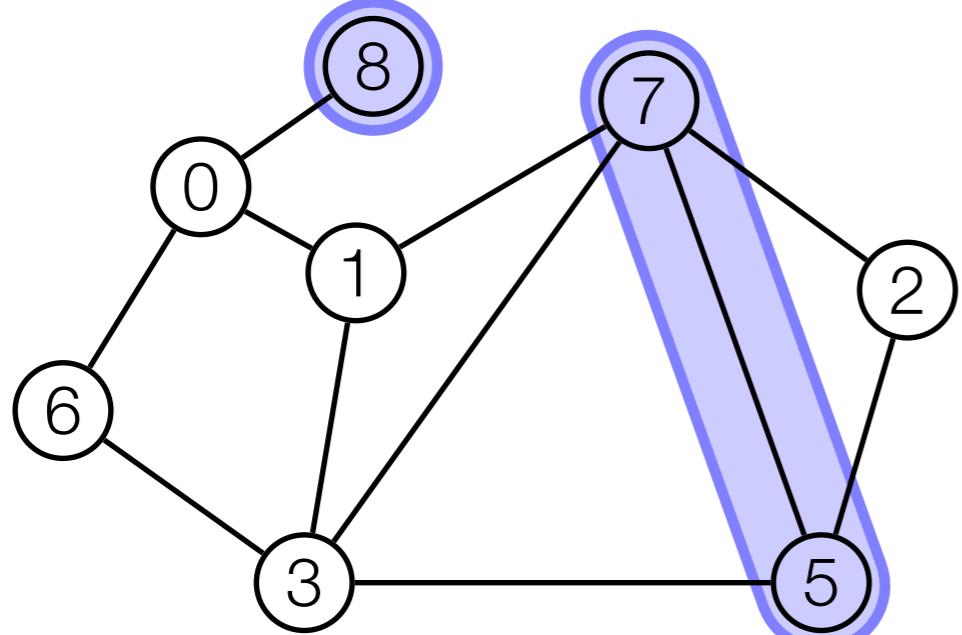
Example: $g = 2, s = 3$



Graph Sparsity Projection

Given a node-weighted graph, find a set of g subtrees with total size s so that the sum of node weights is maximized.

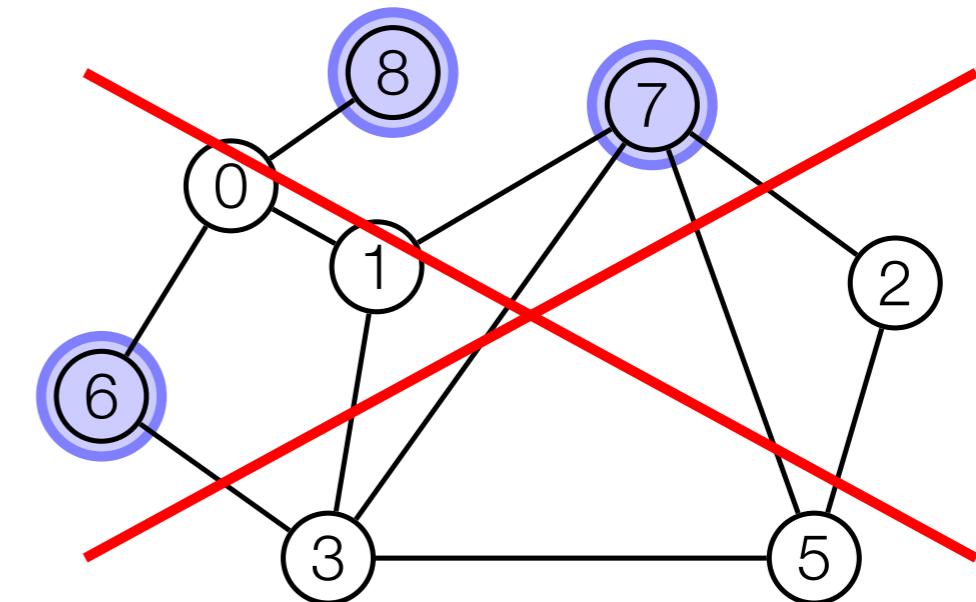
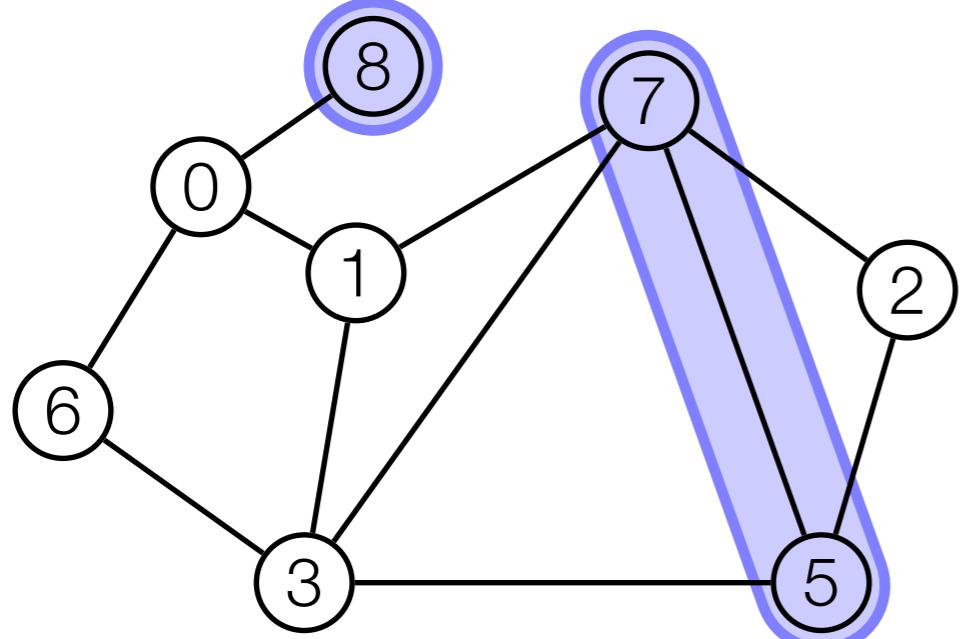
Example: $g = 2, s = 3$



Graph Sparsity Projection

Given a node-weighted graph, find a set of g subtrees with total size s so that the sum of node weights is maximized.

Example: $g = 2, s = 3$



But: This problem is **NP-hard**.

How Can Approximation Help?

Natural relaxation: **“Tail” Approximation**

Given a point $p \in \mathbb{R}^d$, find a $q \in \mathcal{C}$ such that

$$\|p - q\|_2 \leq \min_{q' \in \mathcal{C}} \|p - q'\|_2$$

How Can Approximation Help?

Natural relaxation: **“Tail” Approximation**

Given a point $p \in \mathbb{R}^d$, find a $q \in \mathcal{C}$ such that

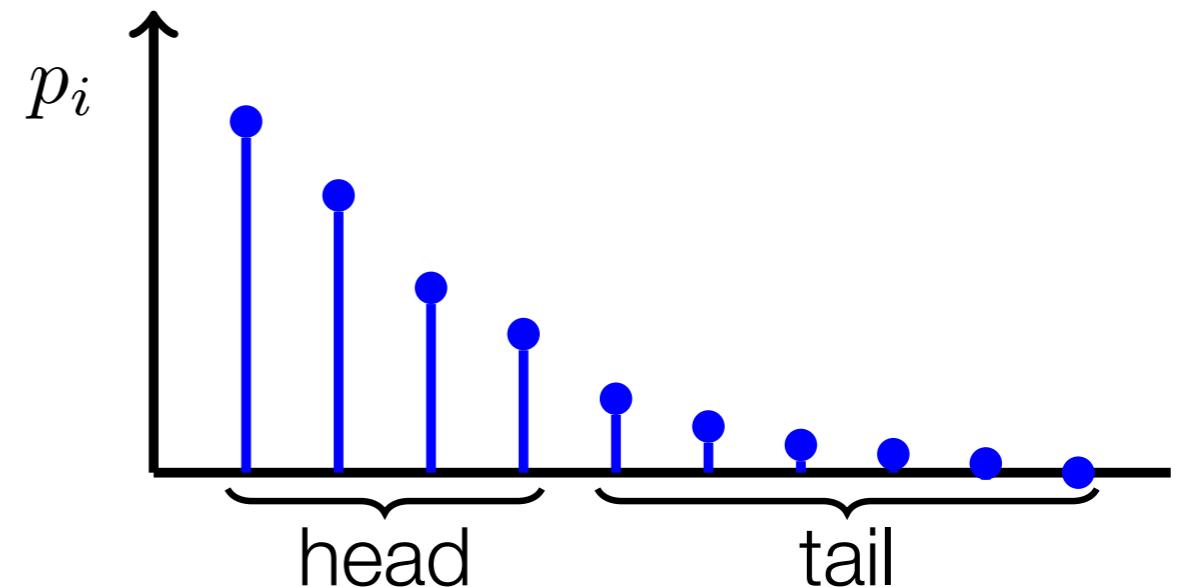
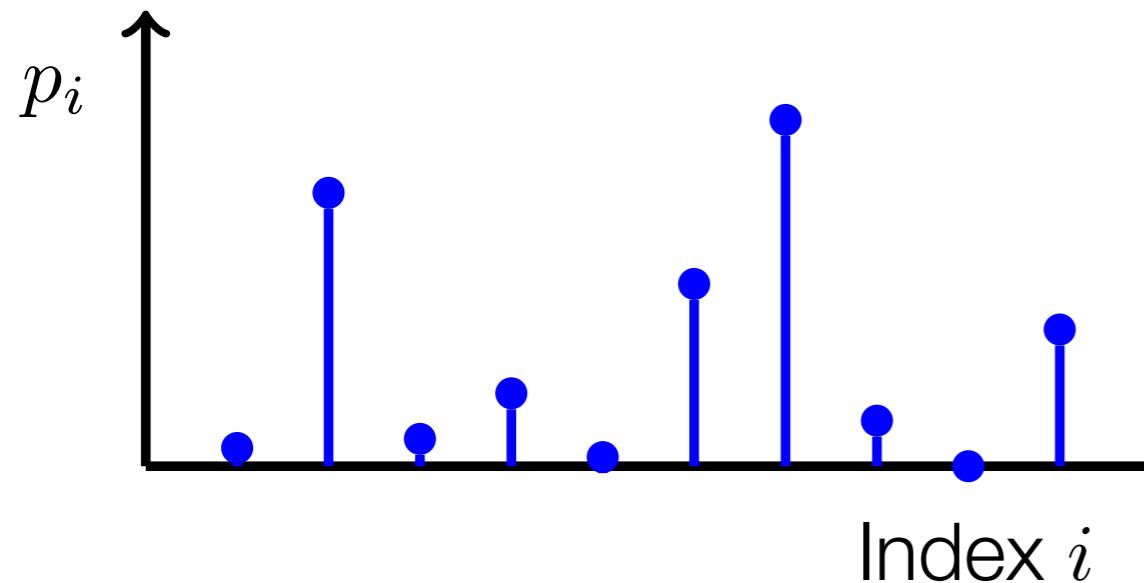
$$\|p - q\|_2 \leq c_{\mathcal{T}} \cdot \min_{q' \in \mathcal{C}} \|p - q'\|_2$$

How Can Approximation Help?

Natural relaxation: “**Tail**” Approximation

Given a point $p \in \mathbb{R}^d$, find a $q \in \mathcal{C}$ such that

$$\|p - q\|_2 \leq c_{\mathcal{T}} \cdot \min_{q' \in \mathcal{C}} \|p - q'\|_2$$



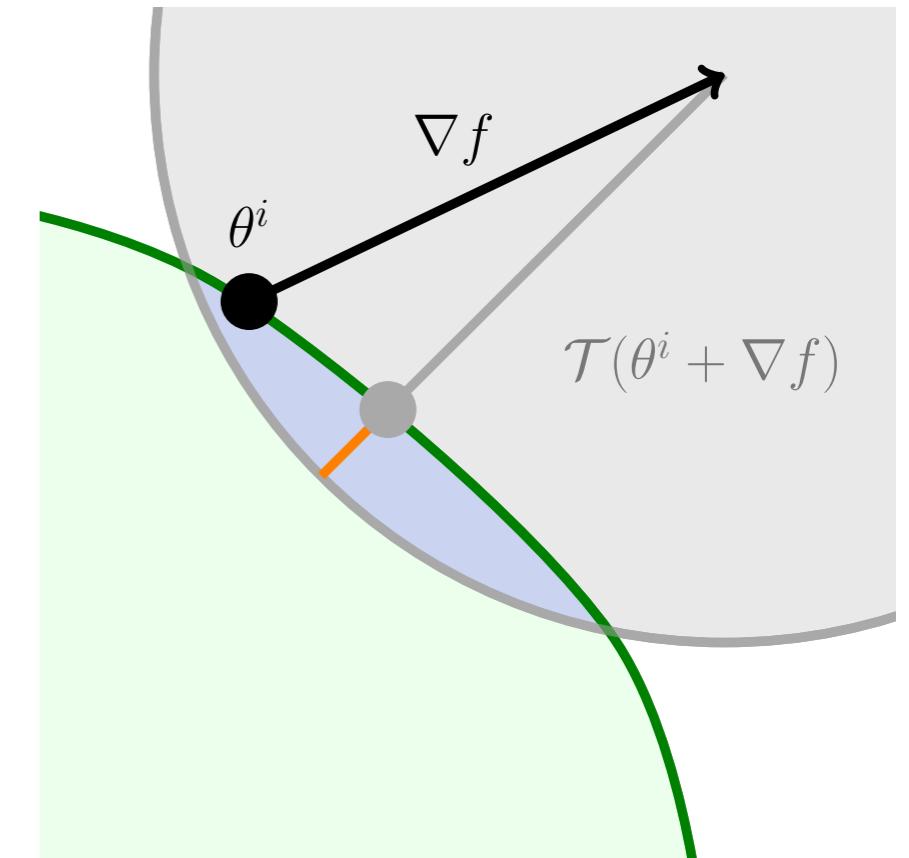
How Can Approximation Help?

Natural relaxation: “**Tail**” Approximation

Given a point $p \in \mathbb{R}^d$, find a $q \in \mathcal{C}$ such that

$$\|p - q\|_2 \leq c_{\mathcal{T}} \cdot \min_{q' \in \mathcal{C}} \|p - q'\|_2$$

**Gradient descent with tail
approximation fails for any $c_{\mathcal{T}} > 1$!**



Simple Fix: More Approximation

Complementary guarantee: **“Head” Approximation**

Given a point $g \in \mathbb{R}^d$, find a $q \in \mathcal{C} \cap \mathbb{S}^{d-1}$ such that

$$\langle g, q \rangle \geq c_{\mathcal{H}} \cdot \max_{q' \in \mathcal{C} \cap \mathbb{S}^{d-1}} \langle g, q' \rangle.$$

Simple Fix: More Approximation

Complementary guarantee: “**Head**” **Approximation**

Given a point $g \in \mathbb{R}^d$, find a $q \in \mathcal{C} \cap \mathbb{S}^{d-1}$ such that

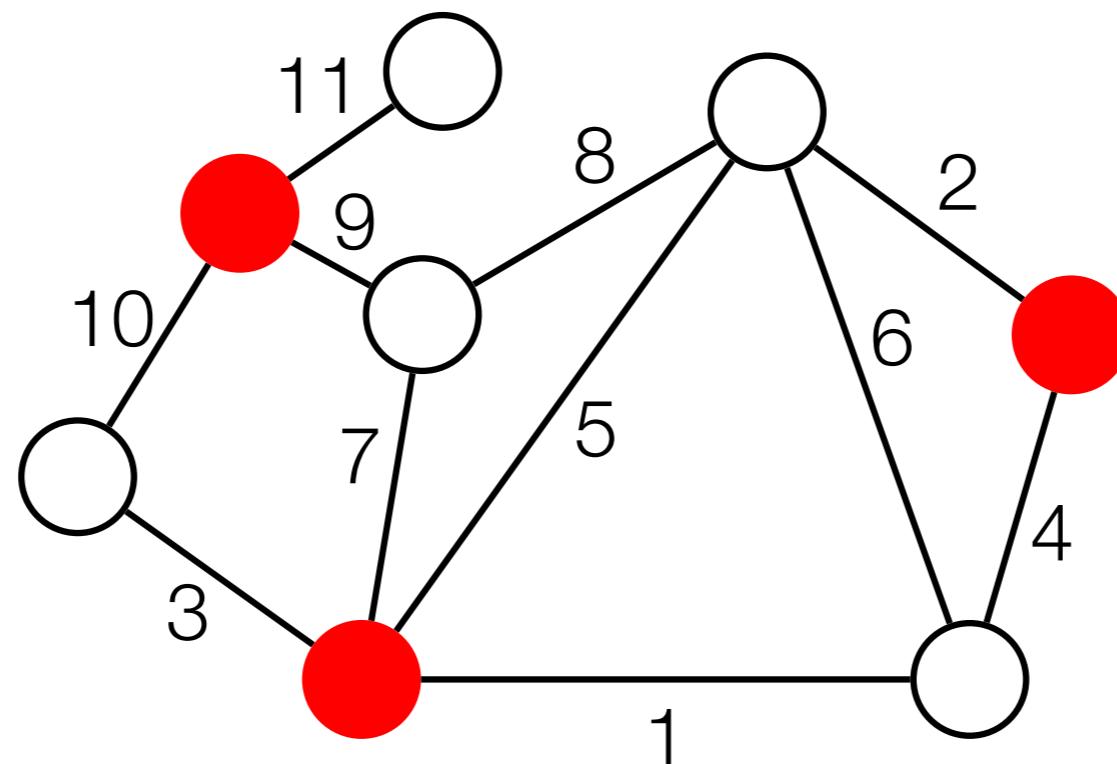
$$\langle g, q \rangle \geq c_{\mathcal{H}} \cdot \max_{q' \in \mathcal{C} \cap \mathbb{S}^{d-1}} \langle g, q' \rangle.$$

Head + Tail approximation combined provably succeeds.

- First analyzed in [Hegde-Indyk-Schmidt’14]
- Generalized in [Cohen-Hegde-Jegelka-Schmidt’??]

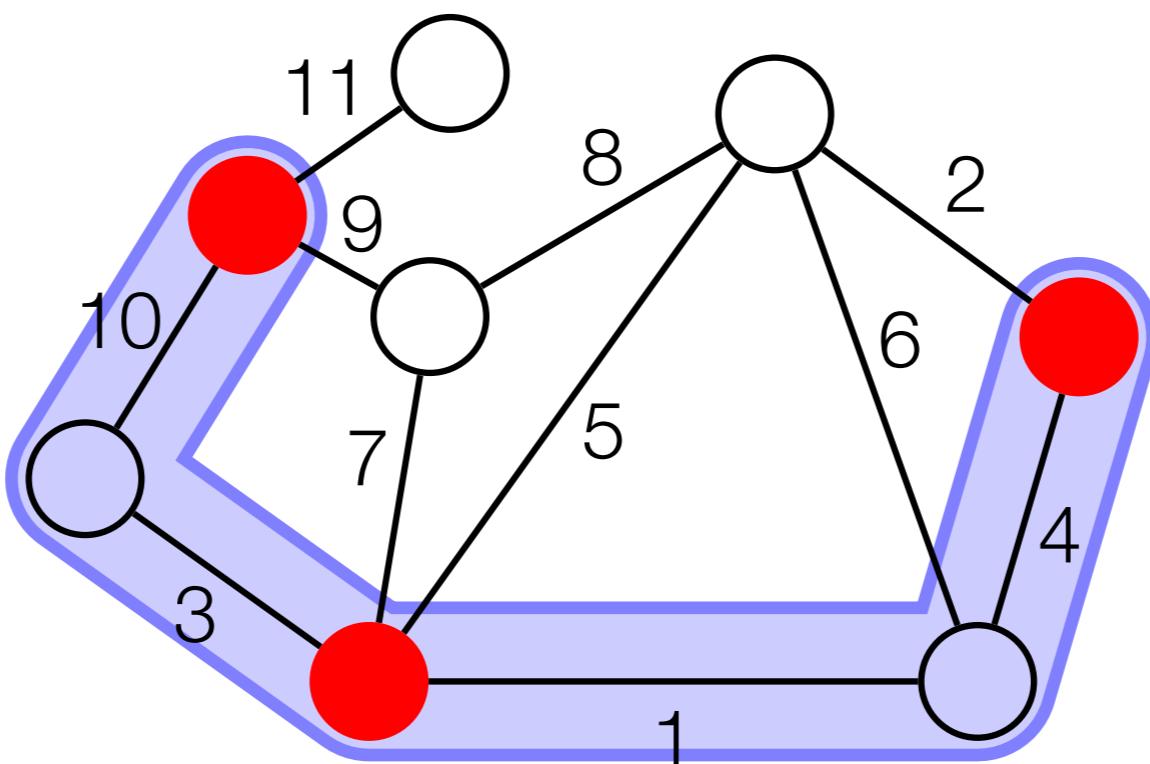
Prize-Collecting Steiner Tree

Generalization of the Steiner tree problem



Prize-Collecting Steiner Tree

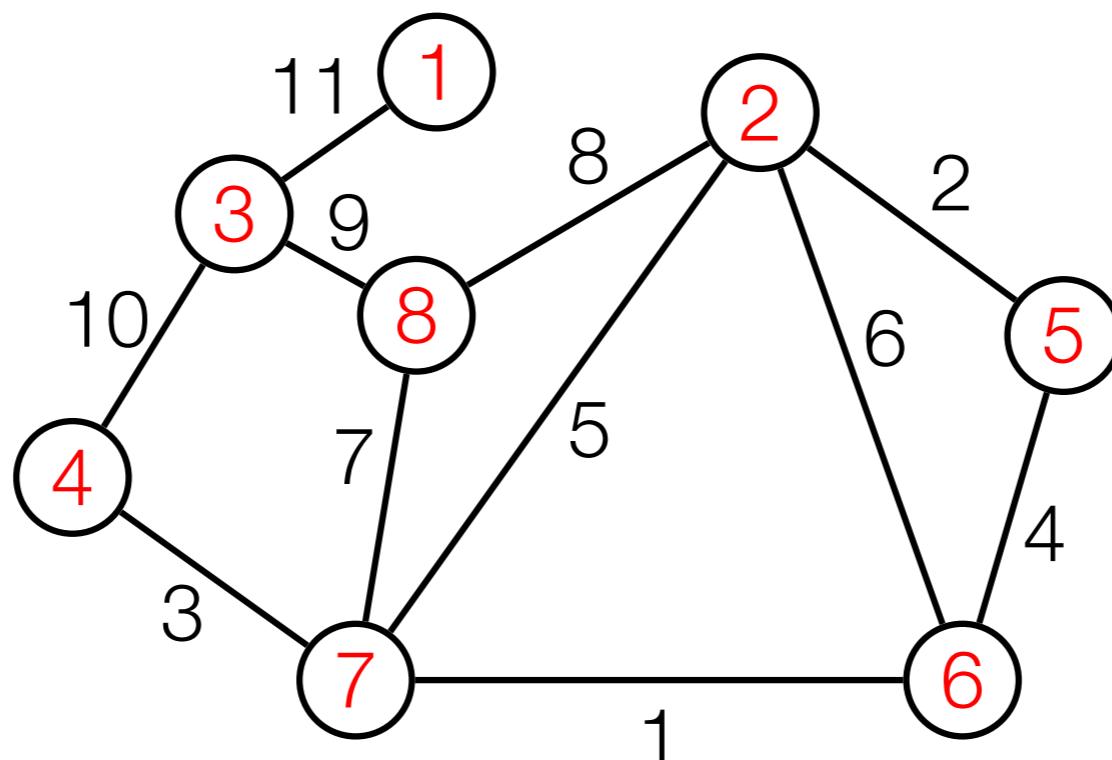
Generalization of the Steiner tree problem



Prize-Collecting Steiner Tree

Generalization of the Steiner tree problem

Every node has a given **prize π** .

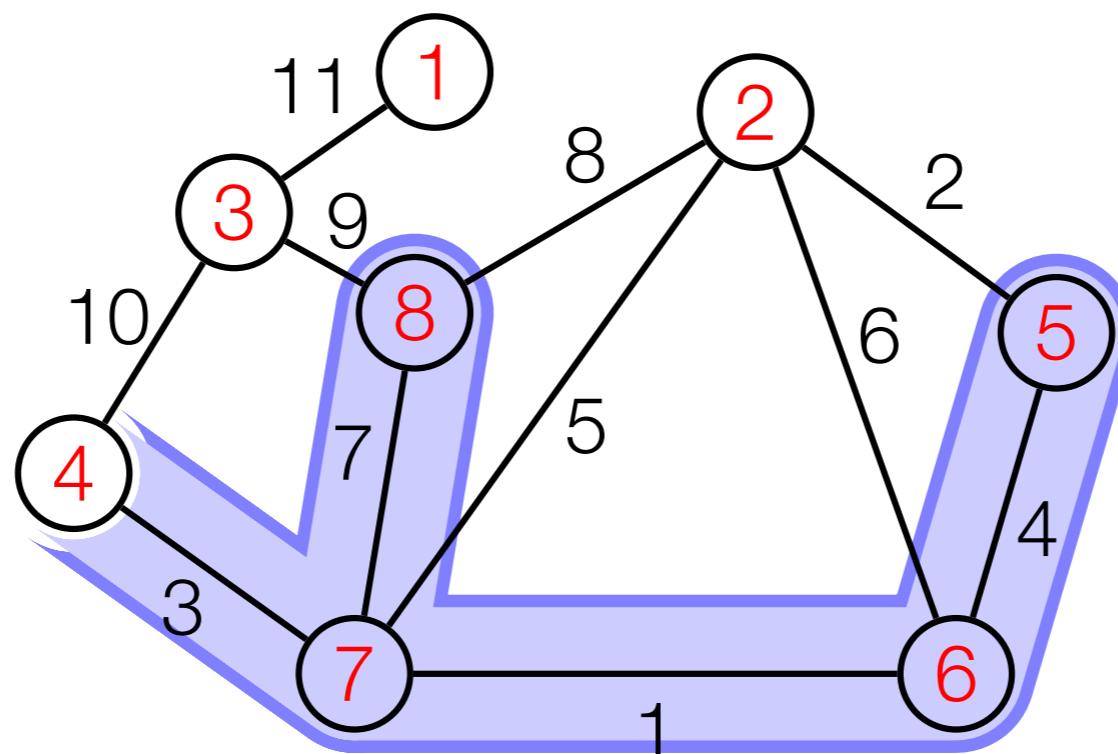


Goal: find a tree T minimizing $\sum_{v \notin V_T} \pi(v) + \sum_{e \in E_T} c(e)$

Prize-Collecting Steiner Tree

Generalization of the Steiner tree problem

Every node has a given **prize π** .

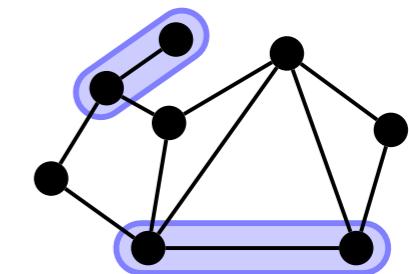


Goal: find a tree T minimizing $\sum_{v \notin V_T} \pi(v) + \sum_{e \in E_T} c(e)$

Connection to Graph Sparsity

Starting point: Goemans-Williamson approximation scheme

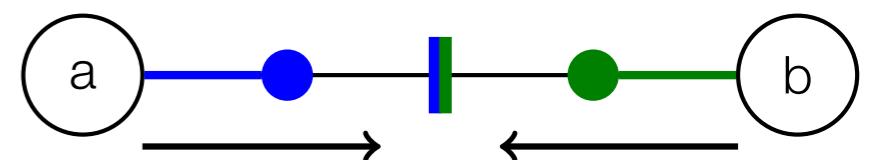
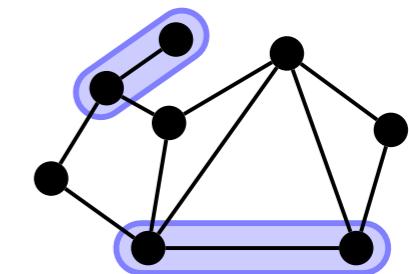
- Generalize to Steiner **forests**



Connection to Graph Sparsity

Starting point: Goemans-Williamson approximation scheme

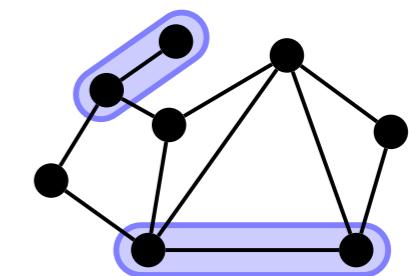
- Generalize to Steiner **forests**
- Develop a **nearly-linear time** and **practical** algorithm



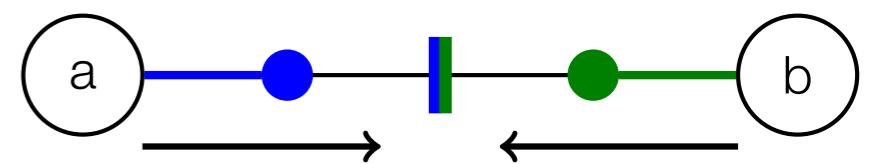
Connection to Graph Sparsity

Starting point: Goemans-Williamson approximation scheme

- Generalize to Steiner **forests**



- Develop a **nearly-linear time** and **practical** algorithm



- **Reduce** graph sparsity projection to a sequence of PCSF instances

$$\sum_{v \notin V_T} \pi(v) + \lambda \cdot \sum_{e \in E_T} c(e)$$

Estimation Guarantee

Setup: $y = X\theta^* + e$

If X is a restricted isometry for graph-sparse vectors,
projected gradient descent converges in $O(\log \|\theta^*\|/\|e\|)$
iterations and finds an estimate $\hat{\theta}$ such that:

$$\|\hat{\theta} - \theta^*\| \leq O(\|e\|)$$

Estimation Guarantee

Setup: $y = X\theta^* + e$

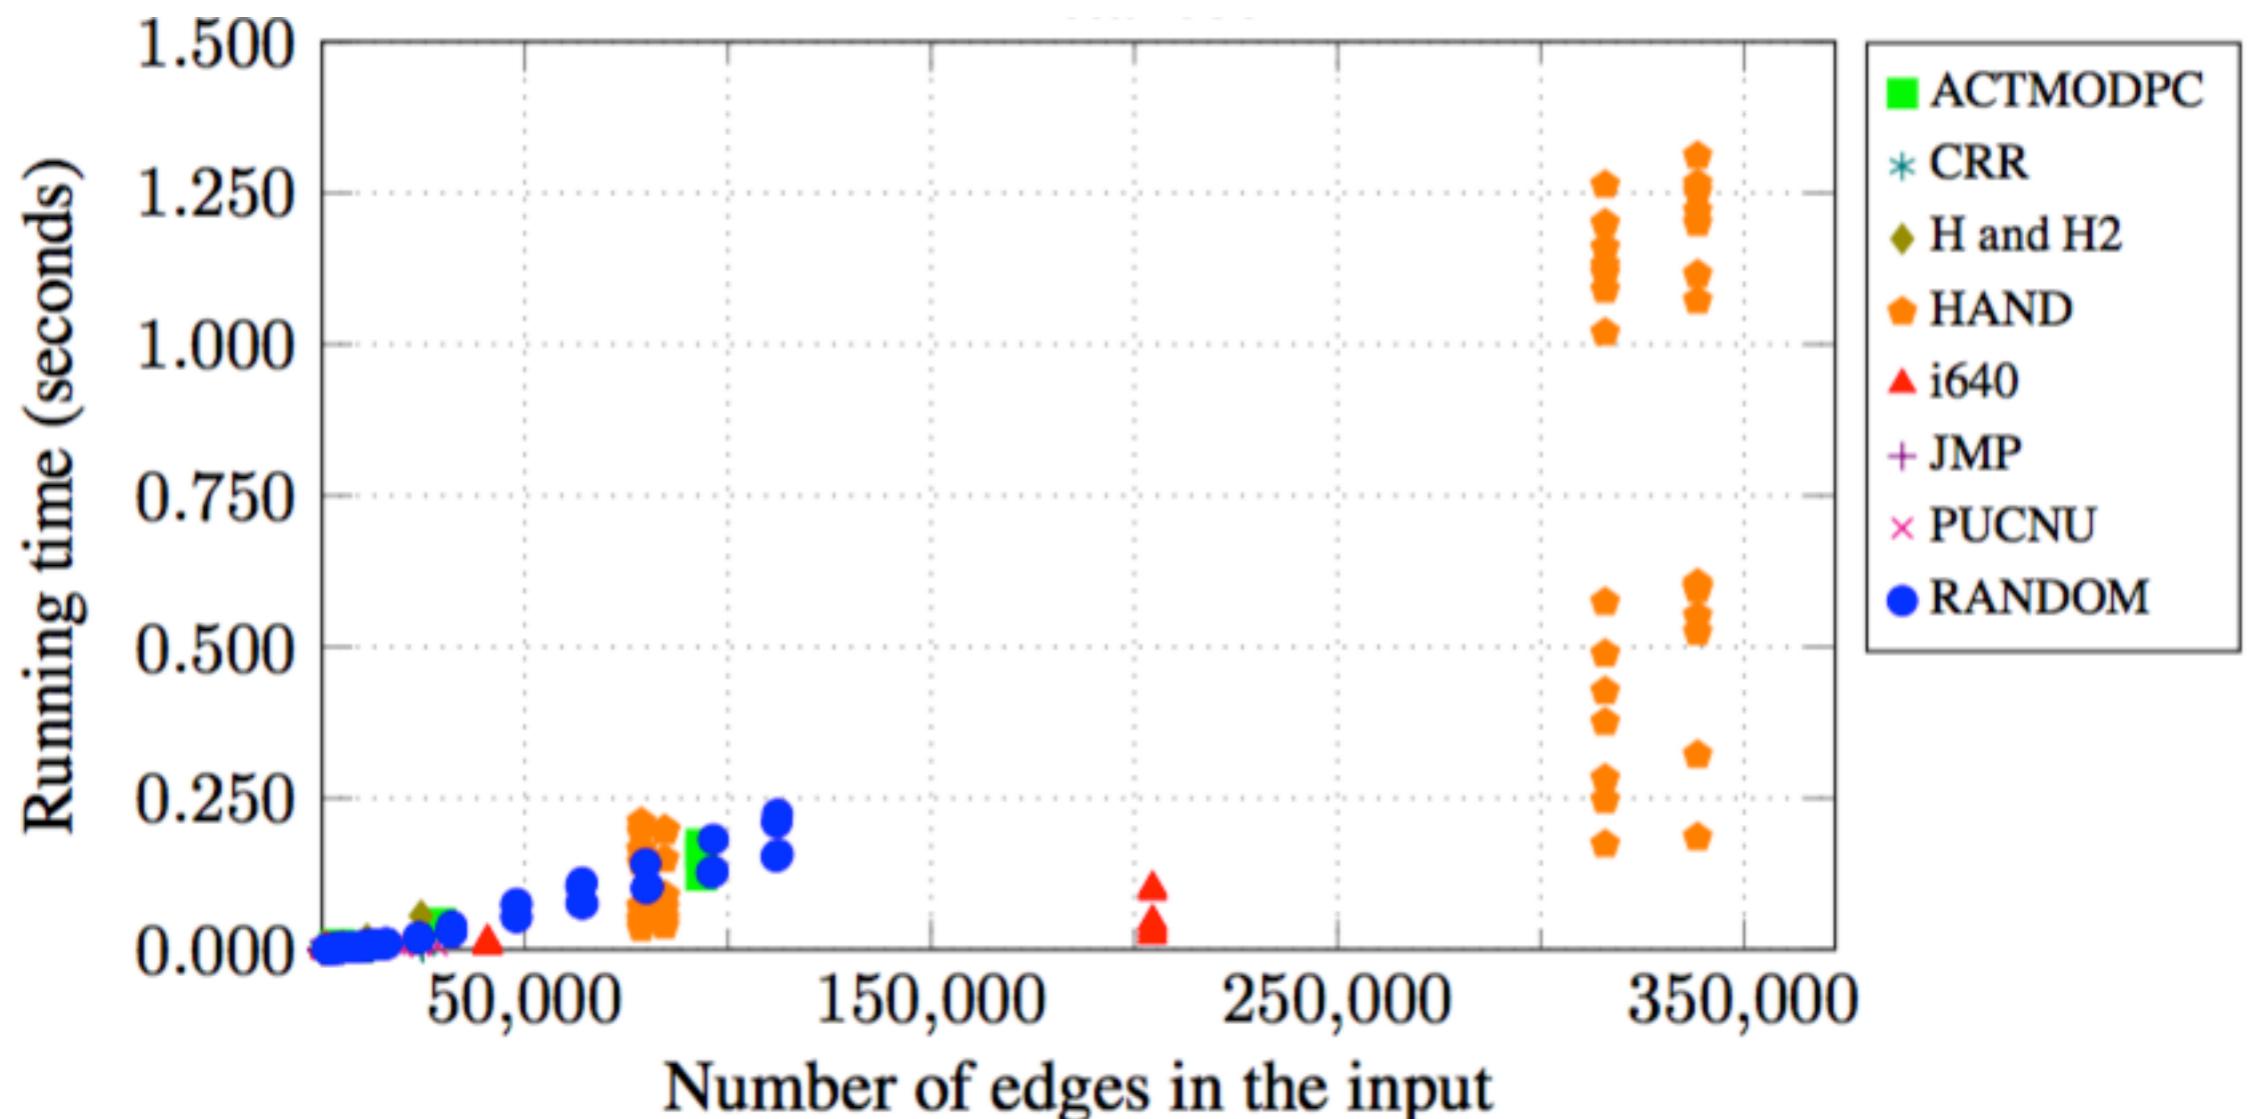
If X is a restricted isometry for graph-sparse vectors, projected gradient descent converges in $O(\log \|\theta^*\|/\|e\|)$ iterations and finds an estimate $\hat{\theta}$ such that:

$$\|\hat{\theta} - \theta^*\| \leq O(\|e\|)$$

Same guarantee for approximate and exact projections.

Experiments: Prize Collecting Steiner Forest

Running time on instance of the 2014 DIMACS challenge.



Nearly-linear scaling with number of edges.

Experiments: Prize collecting Steiner Forest

Algorithm	Solution value	Time
Hegde, Indyk, Schmidt	5601.8	0.2
mozartballs	5058.5	122.1
staynerd	5057.4	143.0
polito	5130.6	160.1
scipjack	5903.5	1052.3
Althaus-Blumenstock	5831.9	2113.5
KTS	5129.0	4842.1
heinz-dc	5695.5	6420.0

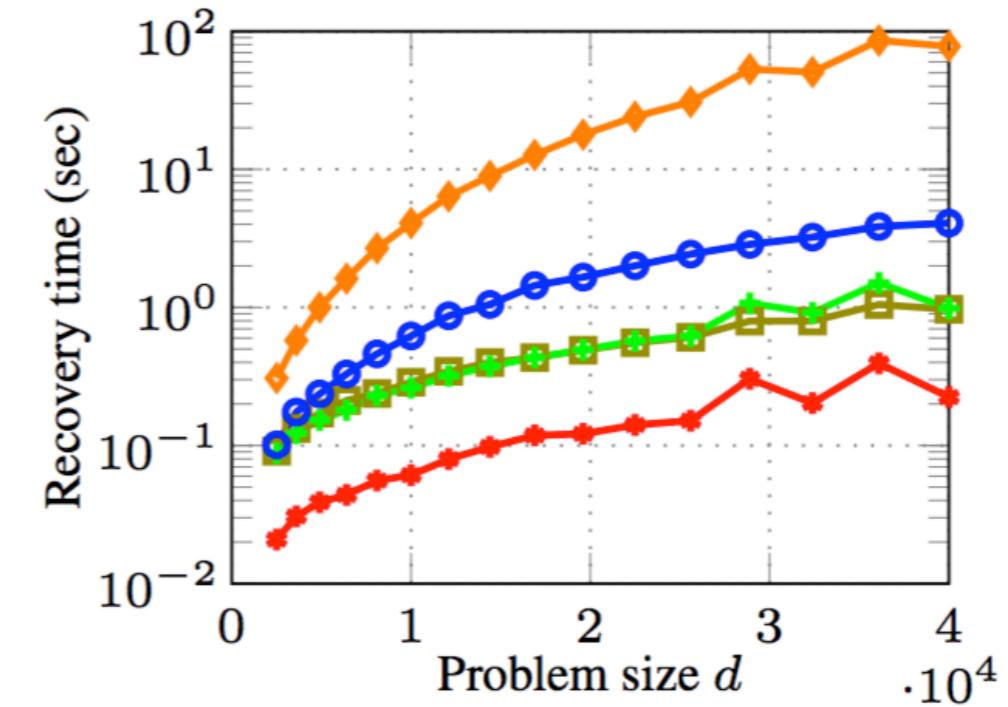
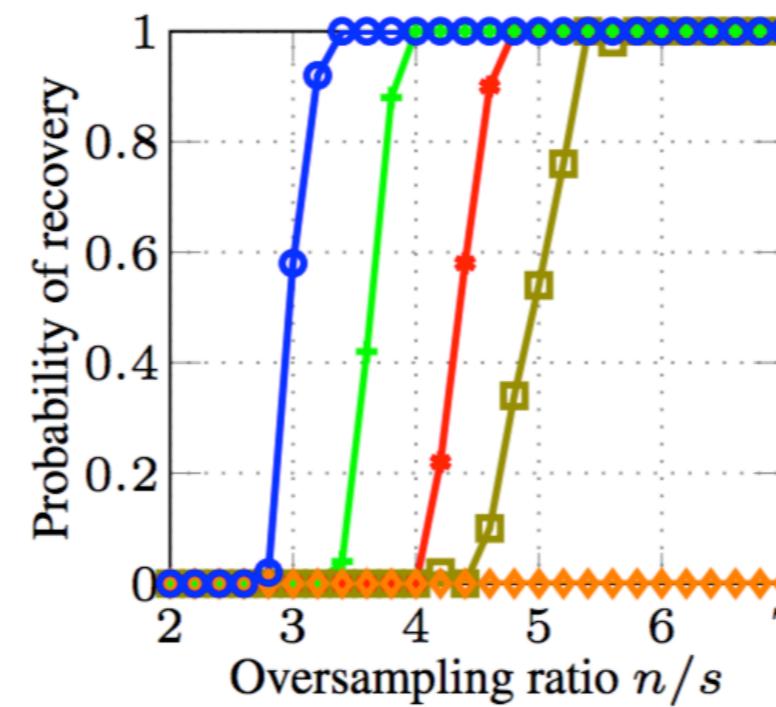
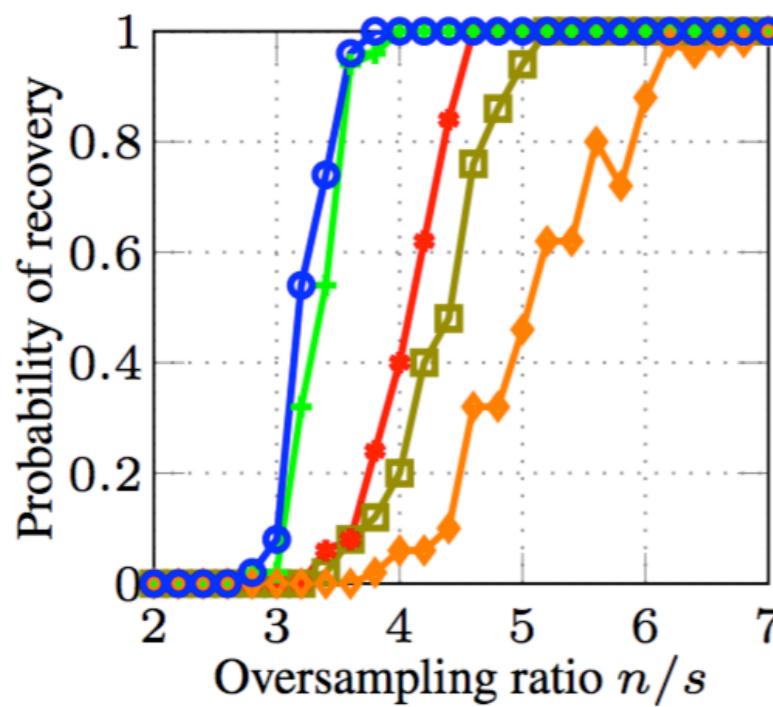
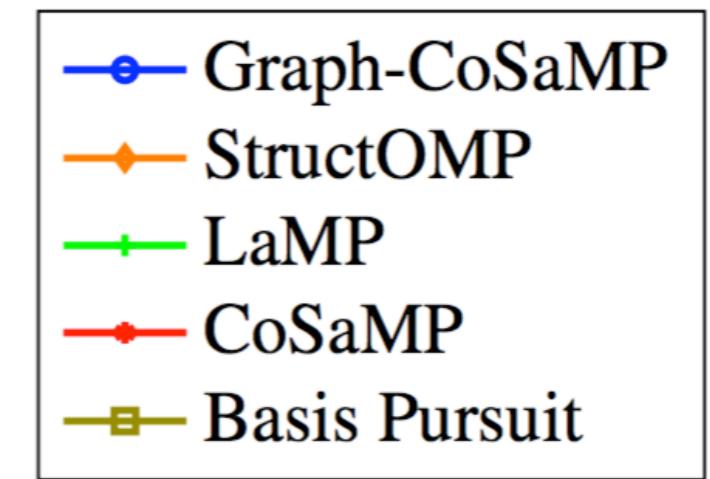
Code for PCSF

The screenshot shows a GitHub repository page for 'fraenkel-lab / pcst_fast'. The repository description is 'A fast implementation of the Goemans-Williamson scheme for the prize-collecting Steiner tree / forest problem.' It has 39 commits, 1 branch, 0 releases, 3 contributors, and is licensed under MIT. The last commit was on Feb 14.

Commit	Message	Date
external	Merge commit '792c14ede5ac83a2577ef9f945baae5648fb5035' as 'external/...'	a year ago
src	many squashed commits which add automatic wheel support	11 months ago
.gitignore	Add a requirements.txt	a year ago
.travis.yml	Update .travis.yml	3 months ago
LICENSE.txt	adding a license file	a year ago

- C++ implementation with Python interface
- Has unit tests
- pip install pcst_fast

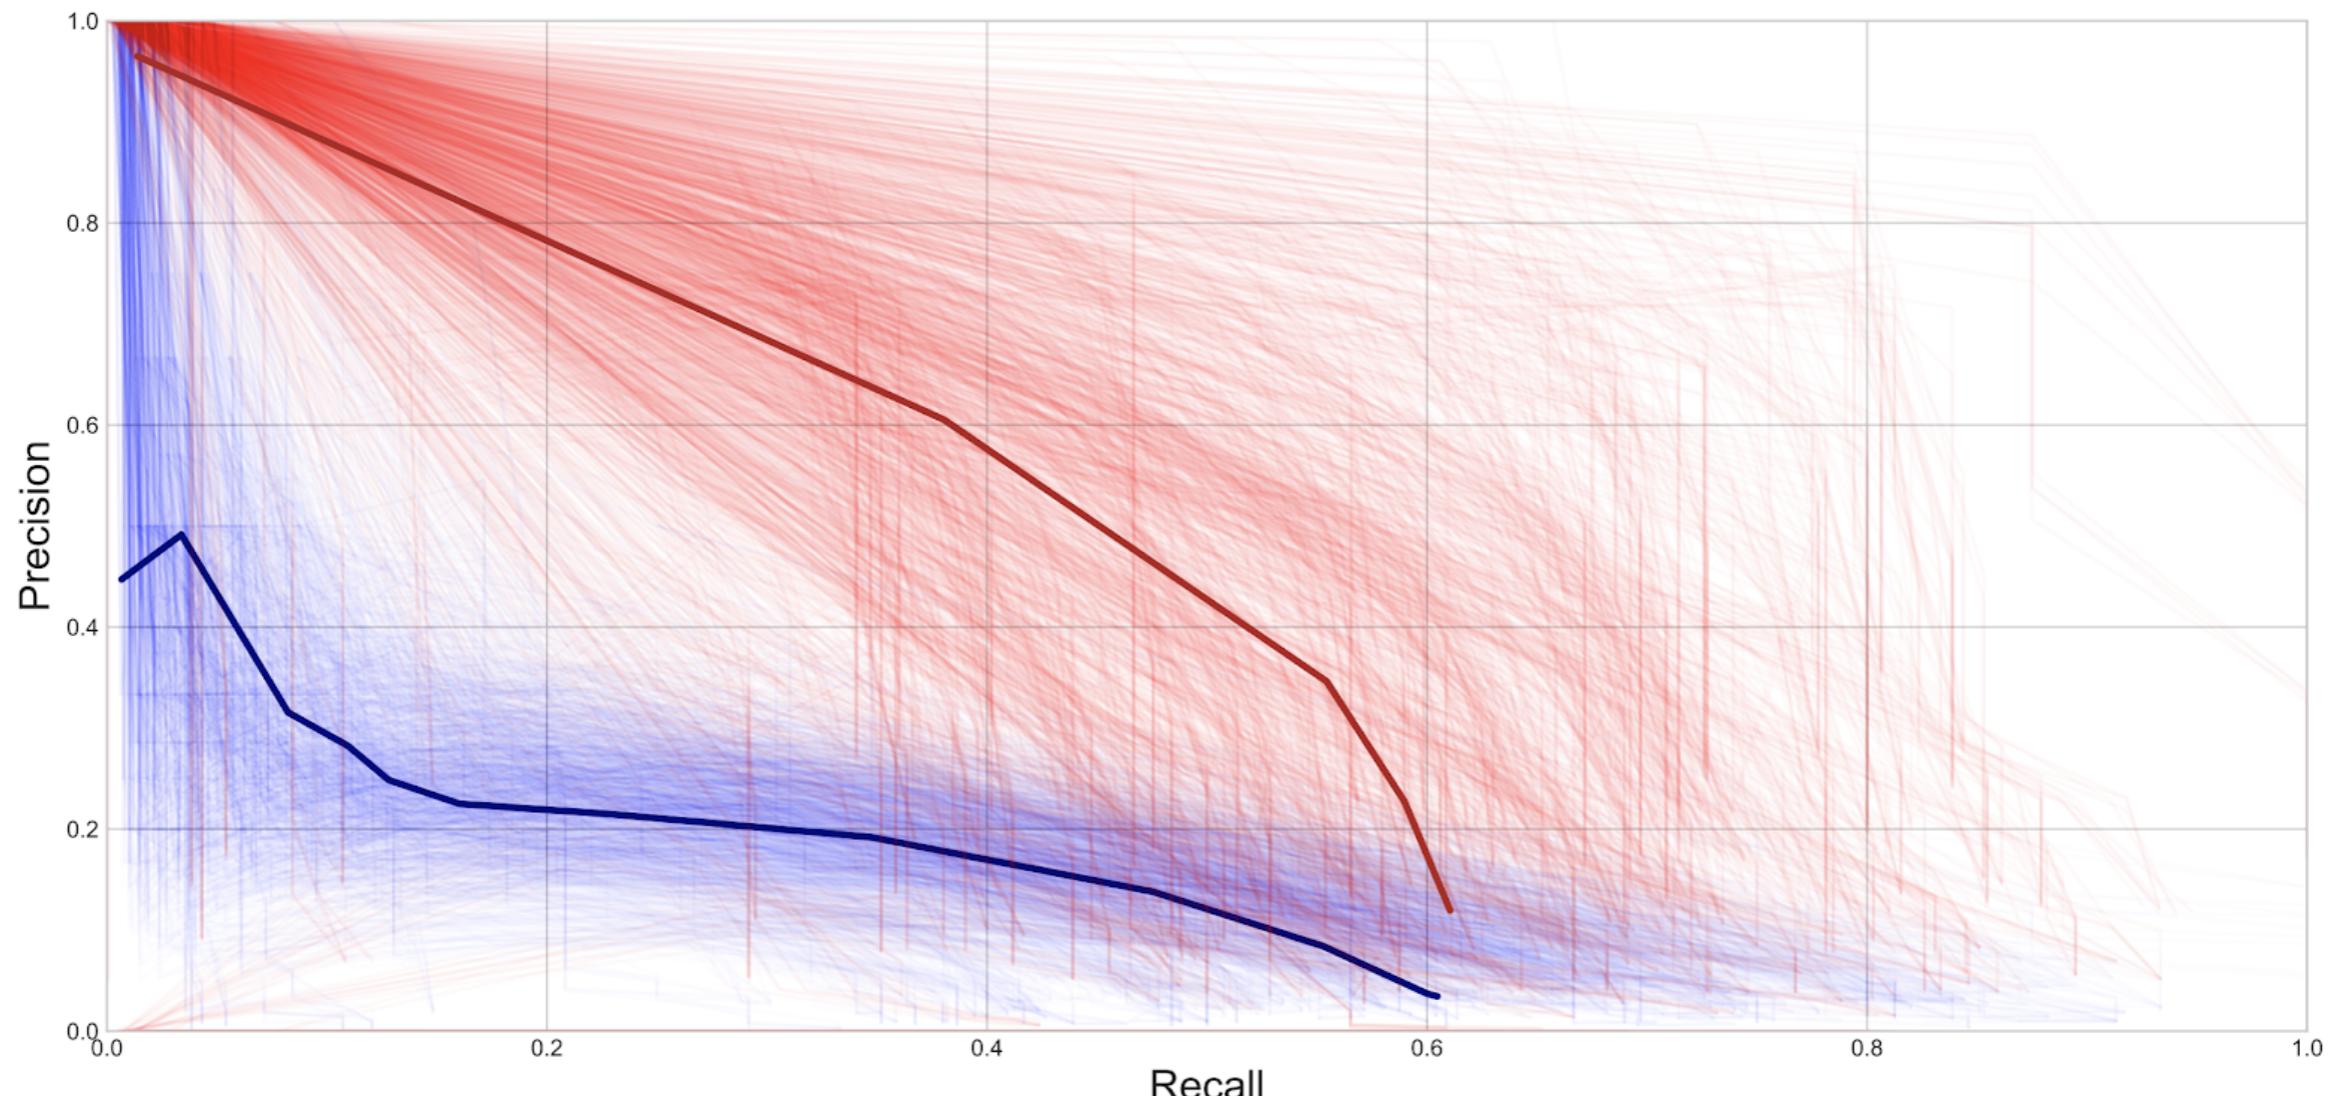
Compressive Sensing Experiments



Graph-Sparse Linear Regression

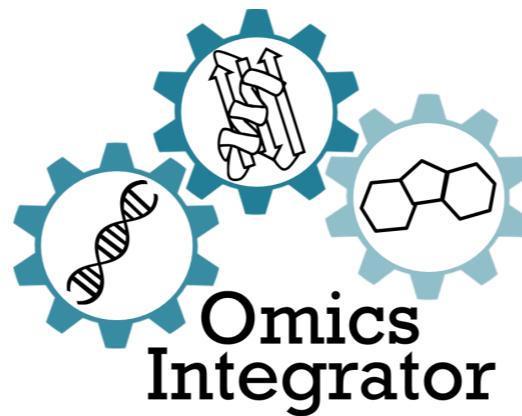
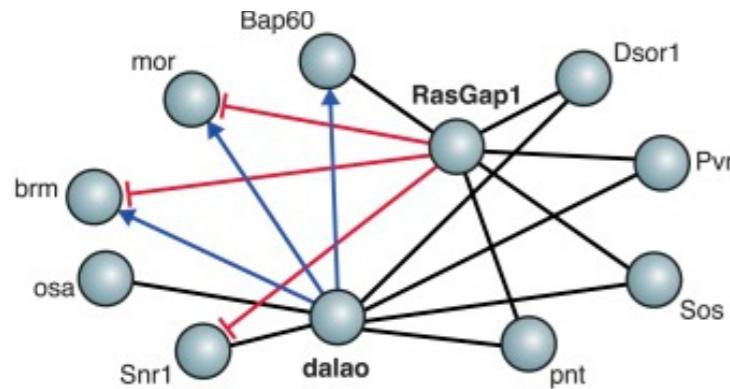
Support recovery on semi-synthetic data.

Graph-sparsity
Lasso



Code on github (but not production-ready yet).

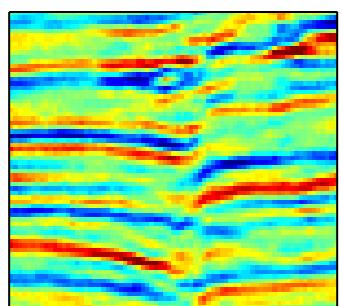
Applications of PCSF



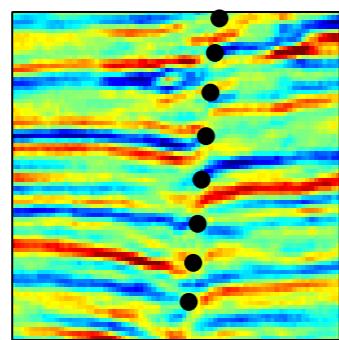
Computational Biology



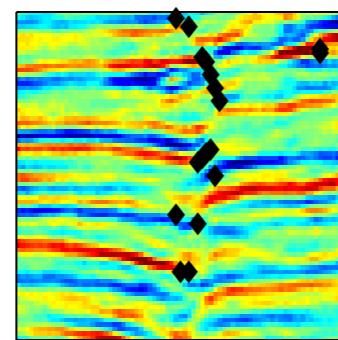
Outage Detection
in Power Networks



Noisy input



Human labels



Automatic

Computational Geophysics

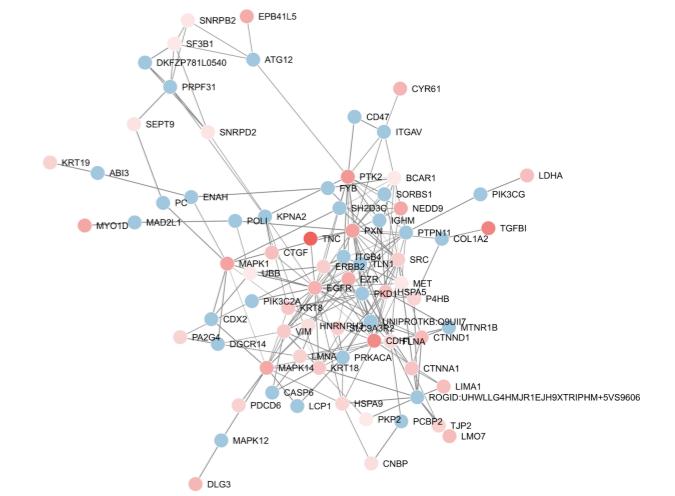


Event Detection
in Social Networks

Conclusions

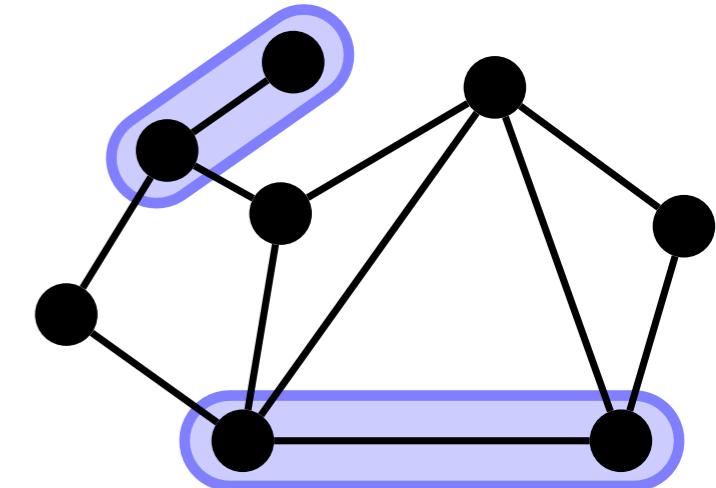
Linear regression with graph constraints

- Find a predictive subgraph.



Fast, rigorous approximations for PCSF

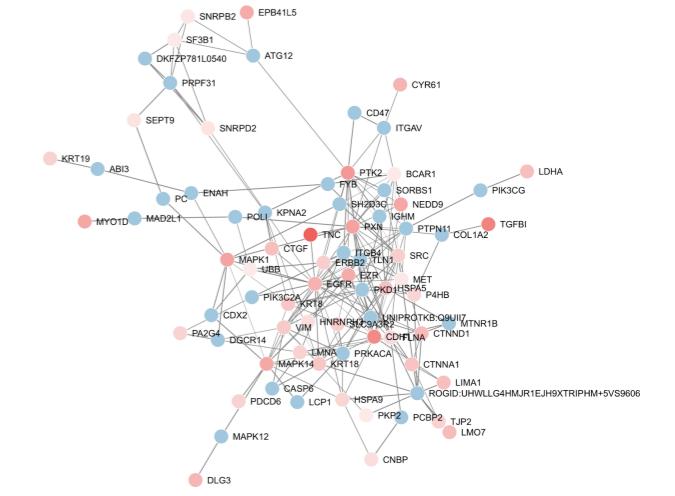
- Collect large node prizes for little edge cost.



Conclusions

Linear regression with graph constraints

- Find a predictive subgraph.



Fast, rigorous approximations for PCSF

- Collect large node prizes for little edge cost.

Future directions

- Applications?
- Other graph constraints in regression?

